



CSI 436/536 (Spring 2026)

Machine Learning

Lecture 8: Loss and Gradient Descent



Chong Liu

Department of Computer Science

Feb 18, 2026

Recap: linear classifier

- Problem of overfitting
 - Too dependent on training data
 - Bad on test data
- Data splitting methods:
 - Holdout
 - Cross validation
- Perceptron algorithm



Pseudo code / implementation for the perceptron algorithm

The Perceptron algorithm

```
Initialize  $w_1 = [0, 0, \dots, 0]$ 
while 1:
    1. Gets feature vector  $x$ 
    2. Agent makes prediction  $\hat{y} \leftarrow \text{sign}(w_t \cdot x)$ 
    3. Reveals label  $y$ 
    4. If  $\hat{y} \neq y$  :
        a.  $w_{t+1} \leftarrow w_t + yx$ 
        b.  $t \leftarrow t + 1$ 
```

Perceptron algorithm makes limited number of mistakes! (if the data is **linearly separable**)

Theorem (Novikoff, 1962): Assume $\|x\|_2 \leq R$ and there exists w^* such that every input satisfies that $\frac{|x \cdot w^*|}{\|w^*\|_2} \geq \gamma$ and $y = \text{sign}(x \cdot w^*)$. Then the **total number of mistakes** Perceptron makes is smaller than R^2/γ^2 .

Remarks:

- The algorithm can run infinitely long.
- No assumption on data distribution.
- Every new data point to predict on is new.
- No hyperparameters to choose. The algorithm does not need to know R, γ .
- Can support infinitely many features!

Using Perceptron over and over

$$\min_{w \in \mathbb{R}^d} \text{Error}(w) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}(h_w(x_i) \neq y_i)$$

- Perceptron can take any sequence of data
- Loop over the data again and again until there is no mistakes.

```
while not converged:  
    converged = true  
    for over data set:  
        make prediction  
        If "Mistake": 1. Update weights. 2. set converged = false
```

- Converge in after at most $\frac{nR^2}{\gamma^2}$ inner loop iterations.

Limitations of the Perceptron Algorithm

- What if the margin is small?
 - The perceptron algorithm will be slower
- What if non-linearly separable?
 - The perceptron algorithm is known to not converge!

Today

- Learn how to train a machine learning classifier!
- Surrogate loss
- Continuous optimization
- Gradient Descent (GD)

Recap: Linear classifier

- Take input feature vector
 - Score(x) = $w_0 + w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4$
 - $x_1 = 1$ (has hyperlinks)
 - $x_2 = 1$ (on contact list)
 - $x_3 = \text{proportion of misspelling}$
 - $x_4 = \text{length}$
- Let label space be $\{-1,1\}$
- Linear classifier:
 - $h_w(x) = \begin{cases} 1, & \text{if Score}(x) \geq 0 \\ -1, & \text{if Score}(x) < 0 \end{cases}$

Key question: How to train linear classifier
(find w)?

$$\min_{w \in \mathbb{R}^d} \text{Error}(w) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}(h_w(x_i) \neq y_i)$$

Discussion:

- 0-1 loss:

$$\mathbf{1}(\text{sign}(w^T x_i) \neq y_i)$$

- Training problem:

$$\min_{w \in \mathbb{R}^d} \text{Error}(w) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}(\text{sign}(w^T x_i) \neq y_i)$$

- How can you minimize 0-1 loss?

0-1 loss is unfortunately very hard to optimize

- 0-1 loss:

$$\mathbf{1}(\text{sign}(w^T x_i) \neq y_i)$$

- Training problem:

$$\min_{w \in \mathbb{R}^d} \text{Error}(w) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}(\text{sign}(w^T x_i) \neq y_i)$$

- Given n data points, the learner needs to check 2^n different configurations.
 - Why 2? Prediction matches / doesn't match label y .
 - It is known as NP-hard.
 - Highly inefficient when n is large.

Just “relax”: relaxing a hard problem into an easier one


$$\min_{w \in \mathbb{R}^d} \text{Error}(w) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}(\text{sign}(w^T x_i) \neq y_i)$$



$$\min_{w \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \ell(w^T x_i, y_i).$$

New loss function is called “surrogate loss”

$$\min_{w \in \mathbb{R}^d} \text{Error}(w) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}(\text{sign}(w^T x_i) \neq y_i)$$



$$\min_{w \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \ell(w^T x_i, y_i).$$

Key point: Choice of surrogate loss must satisfy

$$\mathbf{1}(\text{sign}(w^T x_i) \neq y_i) \leq \ell(w^T x_i, y_i)$$

- Discussion: why?

Loss functions

- 0-1 loss:

$$\mathbf{1}(h_w(x) \neq y) = \mathbf{1}(\text{sign}(S_w(x)) \neq y)$$


- Surrogate losses:

- Logistic loss (binomial deviance):

$$\log_2(1 + \exp(-y \cdot S_w(x)))$$

- Hinge loss (support vector):

$$\max(0, 1 - y \cdot S_w(x))$$



In-class exercise: Intuition of the logistic loss


$$\log_2(1 + \exp(-y \cdot S_w(x)))$$

Try plotting the logistic loss as a function of $y \cdot S_w(x)$


1. What happens when the classifier predicts correctly?
2. What happens when the classifier predicts incorrectly?

Which surrogate loss is easier to minimize?

- Continuous
- Differentiable
 - Except hinge loss, i.e., loss used in “support vector machine (SVM)”
- Convex



Convex vs Nonconvex optimization



- Unique optimum: global/ local.
- Multiple local optima
- In high dimensions possibly exponential local optima

* Be careful: The surrogate loss being convex does not imply all ML problems using surrogate losses are convex. Linear classifiers are, but non-linear classifiers are usually not.


How do we optimize a continuously differentiable function in general?

- The problem:

$$\min_{\theta} f(\theta)$$

- Gradient descent in iterations

$$\theta_{t+1} = \theta_t - \eta_t \nabla f(\theta_t)$$




In-class exercise: gradient descent

- $\min f(x) = x^2$
1. Find x_2 given $x_0 = 2, \eta = 0.1$
 2. Find x_2 given $x_0 = 2, \eta = 0.4$
 3. Find x_2 given $x_0 = 4, \eta = 0.4$
 4. Find x_2 given $x_0 = 2, \eta = 1.5$

Gradient Descent Demo in 2-D

- An excellent demo tool:
 - https://github.com/lilipads/gradient_descent_viz



Gradient of logistic loss for learning a linear classifier

- The function to minimize is

$$\min_{w \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y_i \cdot x_i^T w))$$

- In-class exercise: Calculate the gradient of loss function w.r.t w

$$\nabla f(w) = \frac{1}{n} \sum_{i=1}^n \frac{\exp(-y_i \cdot x_i^T w)}{1 + \exp(-y_i \cdot x_i^T w)} (-y_i x_i)$$

Hint:

- Apply the chain rule.
- $d \log(x) / dx = 1/x$
- $d \exp(x) / dx = \exp(x)$

Drawback: Gradient Descent (GD) uses all data to do one update.

Key question: Is there an efficient way to optimize loss function?