



CSI 436/536 (Spring 2026)

Machine Learning

Lecture 9: Linear Regression

Chong Liu

Department of Computer Science

Feb 23, 2026

Recap: Loss and Gradient Descent

- 0-1 loss in linear classifier

- Hard to optimize!

$$\min_{w \in \mathbb{R}^d} \text{Error}(w) = \frac{1}{n} \sum_{i=1}^n \mathbf{1}(h_w(x_i) \neq y_i)$$

- Surrogate loss

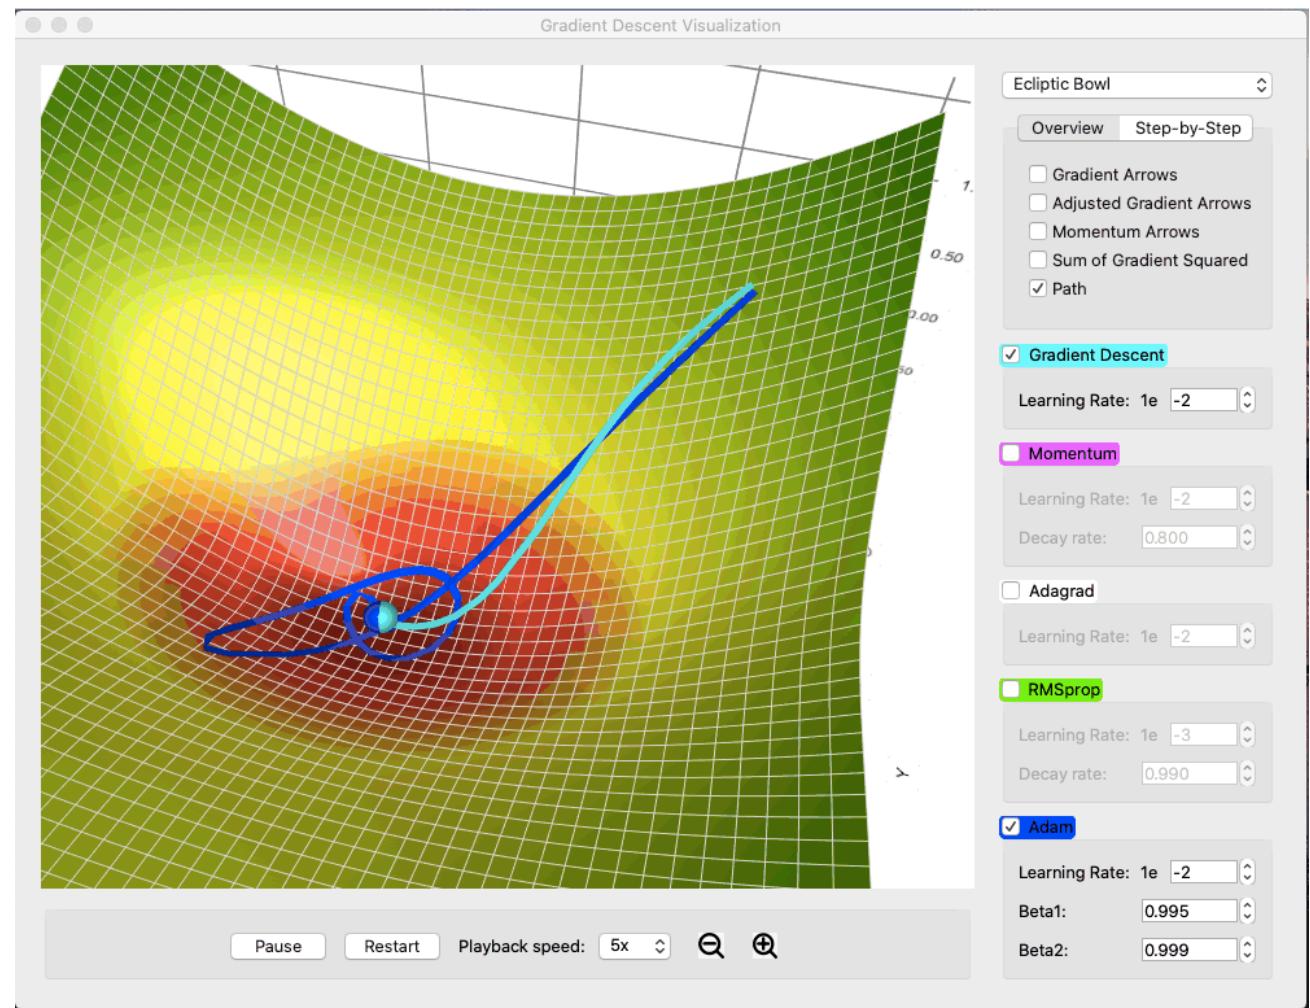
- Easy to optimize (continuous, convex, differentiable)
 - Examples: squared loss, logistic loss, exponential loss, ...

- Gradient Descent (GD)

$$\theta_{t+1} = \theta_t - \eta_t \nabla f(\theta_t)$$

Recap: Gradient Descent Demo in 2-D

- An excellent demo tool:
 - https://github.com/lilipads/gradient_descent_viz

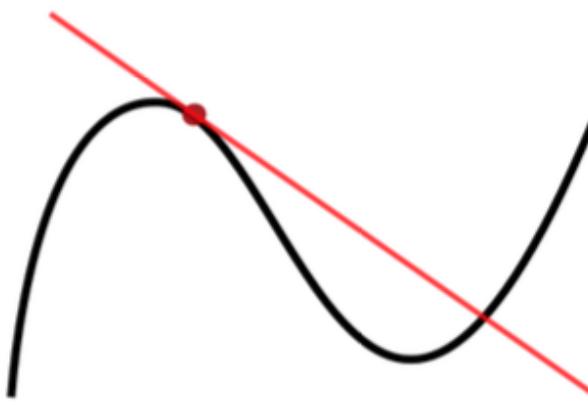


Today

- Stochastic Gradient Descent (SGD)
- Linear regression
- Use SGD to solve linear regression problem!

Recap: What is the “gradient” of a multivariate function?

- We use differentiation to compute derivatives of functions in Calculus.



- Example $f(x, y) = 3x^2 + xy$, $\frac{\partial f(x,y)}{\partial x} = 6x + y$, $\frac{\partial f(x,y)}{\partial y} = x$.
- In many machine learning problems, the objective involves a function that takes a vector of variables as input, e.g., $f(w) = w^T x$ where $w \in \mathbb{R}^d$.
- How to take derivatives on such functions?

Gradient of logistic loss for learning a linear classifier

- The function to minimize is

$$\min_{w \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \log(1 + \exp(-y_i \cdot x_i^T w))$$

- In-class exercise: Calculate the gradient of loss function w.r.t w

$$\nabla f(w) = \frac{1}{n} \sum_{i=1}^n \frac{\exp(-y_i \cdot x_i^T w)}{1 + \exp(-y_i \cdot x_i^T w)} (-y_i x_i)$$

Hint:

- Apply the chain rule.
- $d \log(x) / dx = 1/x$
- $d \exp(x) / dx = \exp(x)$

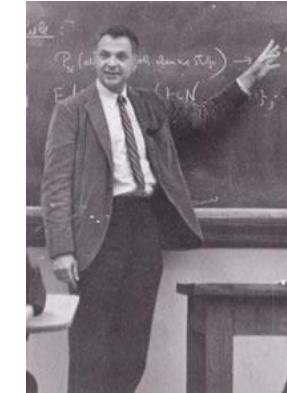
Drawback: Gradient Descent (GD) uses all data to do one update.

Key question: Is there an efficient way to optimize loss function?

Stochastic Gradient Descent (Robbins-Monro 1951)

- Gradient descent

$$\theta_{t+1} = \theta_t - \eta_t \nabla f(\theta_t)$$



Herbert Robbins
1915 - 2001

- Stochastic gradient descent
 - Using a **stochastic approximation** of the gradient:

$$\theta_{t+1} = \theta_t - \eta_t \hat{\nabla} f(\theta_t)$$

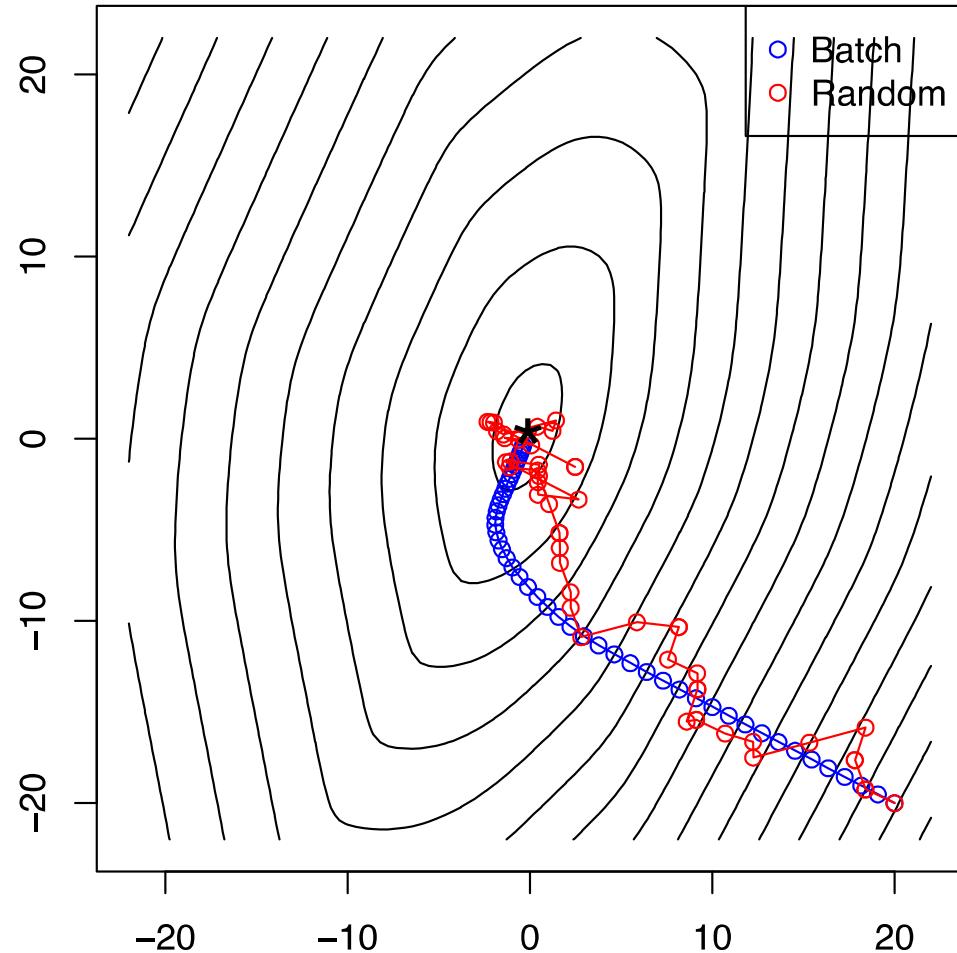
A natural choice of SGD in machine learning

- Recall that

$$\min_{\theta \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \ell(\theta, (x_i, y_i))$$

- SGD samples a data point i uniformly at random while GD uses all data!
 - Use $\nabla_{\theta} \ell(\theta, (x_i, y_i))$

Illustration of **GD** vs **SGD**



Time complexity:

GD: $O(nd * n_{\text{iterations}})$
SGD: $O(d * n_{\text{iterations}})$

Intuition of the SGD algorithm on the “Spam Filter” example

$$\nabla \ell(w, (x_i, y_i)) = \frac{\exp(-y_i \cdot x_i^T w)}{1 + \exp(-y_i \cdot x_i^T w)} (-y_i x_i)$$

Scalar > 0:
≈ 0 if the prediction is correct (no update)
≈ 1 otherwise (update)

Vector of dimension d:
provides the direction of the gradient

Given an email example [1, -1, 0.0375, 80] where 0.0375 is proportion of misspelled words.
Its $y = 1$ means spam.

How will the SGD update change the weight vector?

$$\theta_{t+1} = \theta_t - \eta_t \hat{\nabla} f(\theta_t)$$

If you make a mistake, move the weight towards the direction such that you will be less likely to make the same mistake in the future.

How to choose the step sizes / learning rates?

- In practice:
 - Use cross-validation on validation dataset.
 - Fixed learning rate for SGD is usually fine.
 - If it diverges, decrease the learning rate.

The power of SGD

- Extremely general:
 - Specify an end-to-end differentiable score function
 - E.g., a huge neural network.
- Extremely simple:
 - A few lines of code
- Extremely scalable
 - Just a few pass of the data, no need to store the data

Checkpoint

- Learning a linear classifier:
 - It's hard to directly optimize 0-1 loss
 - Find a surrogate loss
 - Continuous
 - Convex
 - Differentiable
- Gradient descent
 - Calculating gradient / making sense of gradient
 - Improving GD with Stochastic Gradient Descent

Linear regression example: Housing price

- Case study:
 - 8 features:

- MedInc	median income in block group
- HouseAge	median house age in block group
- AveRooms	average number of rooms per household
- AveBedrms	average number of bedrooms per household
- Population	block group population
- AveOccup	average number of household members
- Latitude	block group latitude
- Longitude	block group longitude
 - 1 label: house price
- Discussion: What are they?
 - Feature space (input set)
 - Label space (output set)
 - Linear model
 - Performance metric
 - Loss function

Regression for different problems

- Prediction problem
 - How well can one predict label y ?
 - In housing price example: how well can one predict price given a house?
- Estimation / inference problem
 - How well can one estimate the true function?
 - In housing price example: how well can one learn the price generating function?

Two problems of supervised learning

	Classification		Regression
	Binary classification	Multi-class classification	
Feature space	\mathbb{R}^d	\mathbb{R}^d	\mathbb{R}^d
Label space	$\{-1, 1\}$	$\{1, 2, 3, \dots, K\}$	\mathbb{R}
Performance metric	Classification error (0-1 loss) for test data	Classification error (0-1 loss) for test data	Mean Square Error
Popular surrogate loss (for training)	Logistic loss / exponential loss / square loss	Multiclass logistic loss (Cross-Entropy loss)	Square loss

The objective function for learning linear regression under square loss

- $\hat{w} = \operatorname{argmin}_w \frac{1}{n} \sum_{i=1}^n (x_i^T w - y_i)^2 = \operatorname{argmin}_w \|Xw - y\|_2^2$
 - aka: Ordinary Least Square (OLS)
- In-class exercise: Solve this optimization problem using Direct Solver (setting gradient to 0).

Derive the SGD algorithm

- Problem:
- $\hat{w} = \operatorname{argmin}_w \frac{1}{n} \sum_{i=1}^n (x_i^T w - y_i)^2 = \operatorname{argmin}_w \|Xw - y\|_2^2$
- Step 1: Calculate the gradient of the square loss
- Step 2: Write the SGD update rule

Checkpoint: How to solve linear regression?

- Challenges:
 - We don't have access to future data for prediction!
 - We also don't have access to ground truth
- By solving an optimization problem that **minimizes the loss function on the training data**, and hope that it generalizes.
 - We can verify if it generalizes or not using hold-out / cross-validation ...
- The least square optimization problem using square loss function:

$$\min_{f \in \mathcal{H}} \frac{1}{n} \sum_{i=1}^n (f(x_i) - y_i)^2$$

Checkpoint: Linear regression

- Stochastic Gradient Descent (SGD)
 - Using a **stochastic approximation** of the gradient:
$$\theta_{t+1} = \theta_t - \eta_t \hat{\nabla} f(\theta_t)$$
 - Calculated by one data point randomly sampled from dataset
- Linear regression
 - $\hat{w} = \operatorname{argmin}_w \frac{1}{n} \sum_{i=1}^n (x_i^T w - y_i)^2 = \operatorname{argmin}_w \|Xw - y\|_2^2$
 - Direct solver: $\hat{w} = (X^T X)^{-1} X^T y$
 - GD: $w \leftarrow w - 2\eta X^T (Xw - y)$
 - SGD: $w \leftarrow w - 2\eta x_i (x_i^T w - y_i)$