



# CSI 401 (Fall 2025)

# Numerical Methods

Lecture 19: Course Review

Chong Liu

Department of Computer Science

Nov 24, 2025

# Why learn Numerical Methods?

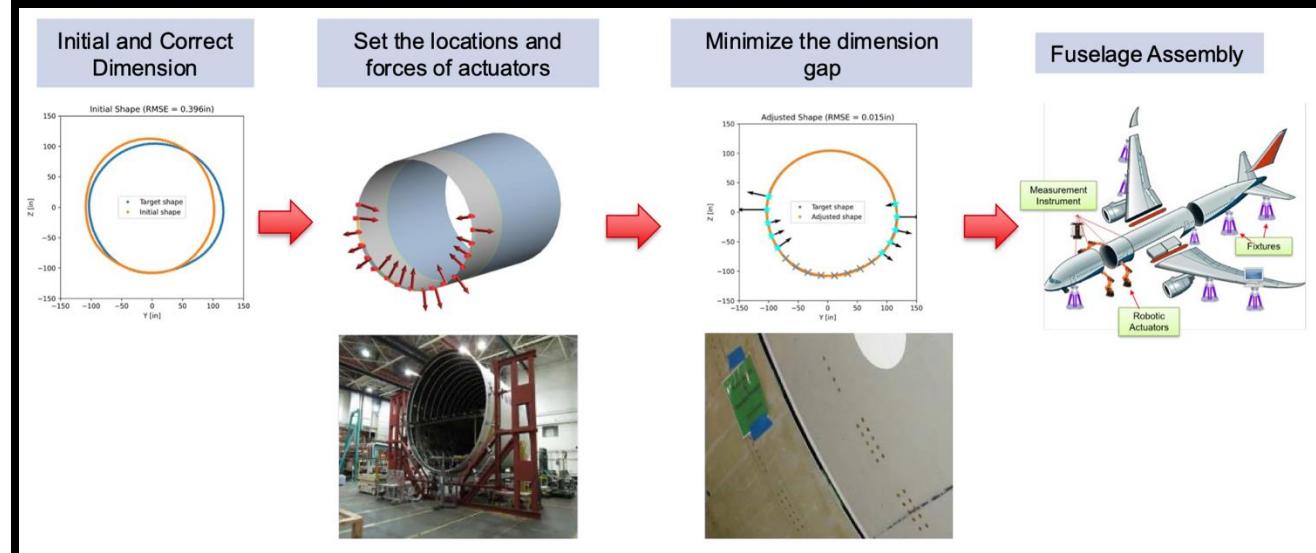
- Motivated by
  - most real-world problems in science, engineering, and data science **cannot** be solved exactly using **closed-form solutions**.
  - For example,
    - Many equations, such as nonlinear systems, high-dimensional integrals, and differential equations, either lack analytical solutions or are too complex to solve by hand.
- Advantages:
  - Providing systematic algorithms to approximate these solutions with controllable accuracy and efficiency
  - Bridging the gap between mathematical theory and computer implementation
  - Enabling the simulation and prediction of complex problems—such as climate modeling, structural design, or machine learning
  - Ensuring stability, convergence, and error control!

# Topics of Numerical Methods covered

- Source of numerical errors
- Asymptotic notations and floating point arithmetic
- Review:
  - Linear algebra, Python, Matlab, LaTeX
- Linear systems:
  - Direct linear equations solvers
  - Eigenvalues and eigenvectors
  - Iterative linear equations solvers
  - Conditioning of linear equations
- Numerical interpolation
  - Data fitting and regression problems
- Nonlinear equations solvers
- Optimization methods
- Numerical integration and differentiation

# Sources of numerical errors

- Data are always noisy
  - Discussion: Any example in your mind?
  - Aircraft fuselage design
- Computers can only handle discrete data
  - Think about data structure: string, array, tree, ...
- No measuring device is perfect
  - Discussion: Any example in your mind?
  - Exact rainfall of Albany, NY in July 2025? No way!



# Types of errors

- Discretization error: we can only deal with values of a function at finitely many points. For example, a very simple way to do numerical differentiation for a function  $f$  is to use a finite difference formula:

$$\hat{f}(x) = \frac{f(x + h) - f(x)}{h}. \quad (2.1)$$

Here, the parameter  $h$  is some small number. It cannot be 0, so this introduces discretization error. Recall that the definition of the derivative of a function at a point  $x$  is

$$\lim_{h \rightarrow 0} \frac{f(x + h) - f(x)}{h}. \quad (2.2)$$

Later in the course, we'll consider better methods than this.

- Convergence error: in which we, say, truncate a power series expansion, stop an iterative algorithm after finitely many iterations, etc.
- Rounding error: This arises because computers have only finite precision. We can only store a finite amount of data in any given machine. Interestingly, in numerical differentiation, there is a tradeoff between discretization error and rounding error (since we cannot make  $h$  infinitely small), and this leads to some optimal choice of  $h$ ! So multiple types of error can play an important role simultaneously in some problems.

# One way is to take the absolute difference

- Definition of **absolute error**:
- $|u - v|$
- It's so simple, but it has some problems:
  - Suppose we try to figure out how much air passes through a Boeing 737 engine per minute during the flight.
  - The true answer is 120,000 pounds.
  - Your solution shows 119,000 pounds, so absolute error is 1,000 pounds.
  - Discussion: You missed 1,000 pounds? Are you doing a good job?



# Another way to measure error

- Definition of relative error:
- $\left| \frac{u-v}{u} \right|$ 
  - note  $u$  is the true number
- Discussion: What's the relative error in previous example?
  - Suppose we try to figure out how much air passes through a Boeing 737 engine per minute during the flight.
  - The true answer is 120,000 pounds.
  - Your solution shows 119,000 pounds.

# Asymptotic notations

- Used to compare the growth of two functions  $f(x)$  and  $g(x)$  as  $x$  tends to some limit point  $x_0$ .
- Discussion:
  - $f(x) = x^2, g(x) = x$ . Which notation shall we use?

To do this, we look at the absolute value of the ratio of the two:

$$\left| \frac{f(x)}{g(x)} \right|. \quad (2.6)$$

The behavior of this can be one of three different things as  $x \rightarrow x_0$ :

•

$$\left| \frac{f(x)}{g(x)} \right| \rightarrow 0. \quad (2.7)$$

In this case, we say that  $f(x)$  is asymptotically negligible compared to  $g(x)$ . We also say that  $f(x) = o(g(x))$  (i.e., “ $f(x)$  is small ‘oh’ of  $g(x)$ ”) as  $x \rightarrow x_0$ .

•

$$\left| \frac{f(x)}{g(x)} \right| \quad (2.8)$$

converges to a positive constant or oscillates but stays bounded.

•

$$\left| \frac{f(x)}{g(x)} \right| \rightarrow \infty. \quad (2.9)$$

In this case, we say that  $f(x)$  is asymptotically dominant compared to  $g(x)$ . This implies that  $g(x) = o(f(x))$ .

# Asymptotic notations

- Additionally, we say that  $f(x) = O(g(x))$  as  $x \rightarrow x_0$  if there is some positive constant  $C$  such that
- $\left| \frac{f(x)}{g(x)} \right| \leq C.$
- Thus, the  $O(\cdot)$  notation means that  $f(x)$  is **asymptotically upper bounded** by  $g(x)$ .
- We also say that  $f(x) = \Omega(g(x))$  if  $g(x) = O(f(x))$ .
- We say that  $f(x) = \Theta(g(x))$  if  $f(x) = O(g(x))$  and  $f(x) = \Omega(g(x))$ .

# Properties of asymptotic notations

**Theorem 3.2** (Properties of asymptotic notations). *Let  $C > 0$  be some positive constant, and let  $x_0 \in \mathbb{R} \cup \{\pm\infty\}$ . Then, for any function  $g(x)$ , as  $x \rightarrow x_0$ ,*

$$C \cdot O(g(x)) = O(g(x)) \quad (3.1)$$

$$C \cdot \Theta(g(x)) = \Theta(g(x)) \quad (3.2)$$

$$C \cdot \Omega(g(x)) = \Omega(g(x)) \quad (3.3)$$

$$C \cdot o(g(x)) = o(g(x)). \quad (3.4)$$

*Additionally, for any  $f(x)$*

$$f(x)O(g(x)) = O(f(x)g(x)). \quad (3.5)$$

- The above theorem allows us to simplify expressions asymptotically. E.g., as  $x \rightarrow \infty$
- $4e^x(\sin(x) + 5) + 3x = \Theta(e^x(\sin(x) + 5)) = \Theta(e^x)$ ,
- where the first equality is because  $x = o(e^x \sin(x))$ , and the second equality is because  $0 \leq |\sin(x)| \leq 1$ .
- Note that all of this can be verified by looking at ratios of functions, as in the definition of the notation.

# Properties of polynomials

**Corollary 3.4.** As  $x \rightarrow \infty$ , for any fixed  $k$ , nonzero constant  $c_k$ , and constants (possibly 0)  $c_j$  for  $j \in \{0, 1, \dots, k - 1\}$ ,

$$P(x) = c_k x^k + c_{k-1} x^{k-1} + \dots + c_1 x + c_0 = \Theta(x^k). \quad (3.8)$$

As  $x \rightarrow 0$ , if  $j$  is the smallest number for which  $c_j \neq 0$ ,

$$P(x) = \Theta(x^j). \quad (3.9)$$

Let us consider the following question: suppose  $f(x) = \Theta(g(x))$ . Is it true in general that  $f(x) - g(x) = \Theta(g(x))$ ? **NO**. For instance,

$$f(x) = 3x, g(x) = 3x + 5 \implies f(x) - g(x) = -5 = o(g(x)). \quad (3.10)$$

# Machine arithmetic - Decimal expansion

- Take 316.1415 for example:

$$316.1415 = 3 \cdot 10^2 + 1 \cdot 10^1 + 6 \cdot 10^0 + 1 \cdot 10^{-1} + 4 \cdot 10^{-2} + 1 \cdot 10^{-3} + 5 \cdot 10^{-4}.$$

- Any real number  $x$  can be written as

$$x = \pm \sum_{j=-\infty}^{\infty} d_j \cdot 10^j$$

- In-class exercise: Decimal expansions for (1) -2, (2)  $\pi$ .

# Machine arithmetic - Binary expansion

- Similar to decimal expansion, every real number  $x$  has a binary (i.e., base  $B = 2$ ) expansion:

$$x = \pm \sum_{j=-\infty}^{\infty} b_j \cdot 2^j$$

- In class exercise: consider a number  $x = -(1011.01)_2$ , what's the binary expansion of it?

$$x = -(1 \cdot 2^3 + 0 \cdot 2^2 + 1 \cdot 2^1 + 1 \cdot 2^0 + 0 \cdot 2^{-1} + 1 \cdot 2^{-2})$$

# Decimal to binary conversion

- Every number has a decimal and a binary expansion. Given a decimal expansion for a number  $x$ , how do we determine its binary expansion?
- We set  $y = x$  and repeatedly do the following:
  - 1. Compute the maximum integer  $j$  such that  $y \geq 2^j$ .
  - 2. Output  $j$ .
  - 3. Compute  $y = y - 2^j$  and go to step 1.
- The algorithm terminates when  $y = 0$ .

# Scientific notation

- Our ultimate goal: come up with a reasonable binary representation of numbers, suitable for storage and manipulation on a computer.
  - Why not just store the binary expansion? The trouble with this is that large numbers can take up a lot more space than smaller numbers, even if they don't have many nonzero digits.
- For instance, consider the following very large number (Avogadro's constant) that arises in chemistry:

$60200000000000000000000000000000$
- How can we store this number in a compact way?

# Scientific notation

Recall how scientific notation works. In decimal, we can write any real number other than 0 as

$$x = \pm m \times 10^E, \quad (5.12)$$

for a unique **mantissa**  $m$  and exponent  $E$ , with  $1 \leq m < 10$  and  $E$  some integer. For example, consider the number 314.159. In scientific notation, this is written as

$$3.14159 \times 10^2. \quad (5.13)$$

In the same fashion, a number can be written in base 2 scientific notation: it takes the form

$$x = \pm m \times 2^E, \quad (5.14)$$

where this time  $1 \leq m < 2$ . For instance, consider the number 3.25. We converted this to binary to get  $(11.01)_2$ . In scientific notation, this becomes

$$(1.101)_2 \times 2^1. \quad (5.15)$$

- In-class exercise: scientific notations of 4125, 40.125, 4.125

# How data are stored? Floating point system

$b_{sign}$	$b_{n_M-1}^{mant} b_{n_M-2}^{mant} \dots b_1^{mant} b_0^{mant}$	$b_{n_E-1}^{exp} b_{n_E-2}^{exp} \dots b_1^{exp} b_0^{exp}$
------------	---	---

Here, there is a single bit giving the sign of the number (0 for negative, 1 for positive). Next is the mantissa, stored as an  $n_M$ -bit number (usually 52 bits). Finally, the exponent is stored as an  $n_E$ -bit number (usually 11 bits). For a nonzero number, the mantissa is not stored directly: since it is between 1 and 2, the binary expansion always begins with a 1. This is redundant, so we **do not** explicitly store it in the floating point representation. It is simply assumed to be there, leading to the so-called **hidden bit representation**. The number 0 has a special representation as all 0s.

- Discussion: what's the number in [0|0100000...|...0000011]?
- Solution:  $-(1+0.25)*8=-10.$

# Linear systems (linear equations)

- An example of linear systems
  - Any linear system can always be rewritten in matrix form

$$a_{11}x_1 + a_{12}x_2 + a_{13}x_3 = b_1$$

$$a_{21}x_1 + a_{22}x_2 + a_{23}x_3 = b_2$$

$$a_{31}x_1 + a_{32}x_2 + a_{33}x_3 = b_3$$

$$\begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix} = \begin{bmatrix} b_1 \\ b_2 \\ b_3 \end{bmatrix}$$

- More generally,  $Ax = b$ 
  - $A$  is an  $m \times n$  matrix
  - $x$  is an  $n$ -dimensional vector
  - $b$  is an  $m$ -dimensional vector
- Problem: given  $A$  and  $b$ , how can you solve  $x$ ?

# Gaussian elimination

- Rewrite the problem in augmented matrix form
  - Original augmented matrix and manipulated augmented matrix

$$\begin{bmatrix} 1 & 1 & 1 & 7 \\ 3 & 2 & 1 & 11 \\ 4 & -2 & 2 & 8 \end{bmatrix} \xrightarrow{\hspace{1cm}} \begin{bmatrix} 1 & 1 & 1 & 7 \\ 0 & -1 & -2 & -10 \\ 0 & 0 & 10 & 40 \end{bmatrix}$$

- Elementary row operations:
  - $2^{\text{nd}}$  line =  $2^{\text{nd}}$  line –  $3 * 1^{\text{st}}$  line [0 -1 -2 -10], done!
  - $3^{\text{rd}}$  line =  $3^{\text{rd}}$  line –  $4 * 1^{\text{st}}$  line [0 -6 -2 -20]
  - Discussion: how to proceed?
  - $3^{\text{rd}}$  line =  $3^{\text{rd}}$  line –  $6 * 2^{\text{nd}}$  line [0 0 10 40], done!

# Gaussian elimination

- Key idea:
  - Use elementary row operations to make A become a right-triangular matrix
  - So that you can sequentially solve the linear systems bottom-up!

$$\left[ \begin{array}{cccc|c} a'_{11} & a'_{12} & \cdots & a'_{1k} & b'_1 \\ 0 & a'_{22} & \cdots & a'_{2k} & b'_2 \\ \vdots & \vdots & \ddots & \vdots & \vdots \\ 0 & 0 & \cdots & a'_{kk} & b'_k \end{array} \right].$$

# Gauss-Jordan Elimination: Beyond Gaussian Elimination

- Consider this linear system:
$$\begin{aligned}x - 2y + 3z &= 9 \\ -x + 3y &= -4 \\ 2x - 5y + 5z &= 17\end{aligned}$$

- Yes! It's Gauss-Jordan Elimination.

- Key idea:
  - Use elementary row operations to make A become an identity matrix
  - So that you can directly read the results!

$$\begin{array}{cccc|c} 1 & -2 & 3 & 9 \\ -1 & 3 & 0 & -4 \\ 2 & -5 & 5 & 17 \end{array} \xrightarrow[\text{Gaussian elimination}]{} \begin{array}{cccc|c} 1 & -2 & 3 & 9 \\ 0 & 1 & 3 & 5 \\ 0 & 0 & 1 & 2 \end{array} \xrightarrow[\text{Elementary row operations}]{} \begin{array}{cccc|c} 1 & 0 & 0 & 1 \\ 0 & 1 & 0 & -1 \\ 0 & 0 & 1 & 2 \end{array}$$

# LU decomposition

- $A = LU$ 
  - $L$  is a lower triangular matrix,  $U$  is a upper triangular matrix
  - Note this decomposition may **not** be unique
- In-class exercise:
  - Find an LU decomposition of  $A = \begin{bmatrix} 2 & 1 & 1 \\ 4 & -6 & 0 \\ -2 & 7 & 2 \end{bmatrix}$ 
    - $u_{11} = a_{11} = 2, u_{12} = a_{12} = 1, u_{13} = a_{13} = 1$
    - $\ell_{21} = a_{21}/u_{11} = 4/2 = 2, \ell_{31} = a_{31}/u_{11} = -2/2 = -1$
    - $u_{22} = a_{22} - \ell_{21}u_{12} = -6 - 2 \cdot 1 = -8$
    - $u_{23} = a_{23} - \ell_{21}u_{13} = 0 - 2 \cdot 1 = -2$
    - $\ell_{32} = (a_{32} - \ell_{31}u_{12})/u_{22} = (7 - (-1) \cdot 1)/(-8) = 8/(-8) = -1$
    - $u_{33} = a_{33} - \ell_{31}u_{13} - \ell_{32}u_{23} = 2 - (-1) \cdot 1 - (-1) \cdot (-2) = 2 + 1 - 2 = 1$
- Solutions:
$$L = \begin{bmatrix} 1 & 0 & 0 \\ 2 & 1 & 0 \\ -1 & -1 & 1 \end{bmatrix}, \quad U = \begin{bmatrix} 2 & 1 & 1 \\ 0 & -8 & -2 \\ 0 & 0 & 1 \end{bmatrix}.$$

# Partial pivoting prevents this issue

- How does partial pivoting work?
  - Swap rows to make pivot have the largest absolute value in its column.

$$\left( \begin{array}{cc|c} 10^{-20} & 1 & 1 \\ 1 & 1 & 2 \end{array} \right) \xrightarrow{R_1 \leftrightarrow R_2} \left( \begin{array}{cc|c} 1 & 1 & 2 \\ 10^{-20} & 1 & 1 \end{array} \right) \xrightarrow{R_2 \leftarrow R_2 - 10^{-20}R_1} \left( \begin{array}{cc|c} 1 & 1 & 2 \\ 0 & 1 - 10^{-20} & 1 - 2 \cdot 10^{-20} \end{array} \right)$$
$$x_2 = \frac{1 - 2 \cdot 10^{-20}}{1 - 10^{-20}} \approx 1 \pm 10^{-20}$$
$$x_1 + x_2 = 2 \implies x_1 = 1 \pm 10^{-20}.$$

- Why does it work?
  - Avoid dividing by tiny numbers, reduces relative error, and makes LU numerically stable for most matrices.

# Another example of partial pivoting

$$\begin{pmatrix} 2 & 1 & 3 \\ 1 & 10^{-20} & 4 \\ 7 & -20 & 5 \end{pmatrix} \xrightarrow{R_1 \leftrightarrow R_3} \begin{pmatrix} 7 & -20 & 5 \\ 1 & 10^{-20} & 4 \\ 2 & 1 & 3 \end{pmatrix} \xrightarrow{R_2 \leftarrow R_2 - (1/7)R_1, R_3 \leftarrow R_3 - (2/7)R_1} \begin{pmatrix} 7 & -20 & 5 \\ 0 & 2.857 & 3.286 \\ 0 & 6.714 & 1.5714 \end{pmatrix} \quad (11.18)$$

$$\xrightarrow{R_2 \leftrightarrow R_3} \begin{pmatrix} 7 & -20 & 5 \\ 0 & 6.714 & 1.5714 \\ 0 & 2.857 & 3.286 \end{pmatrix} \xrightarrow{R_3 \leftarrow R_3 - (2.857/6.714)R_2} \begin{pmatrix} 7 & -20 & 5 \\ 0 & 6.714 & 1.5714 \\ 0 & 0 & 2.6173 \end{pmatrix} \quad (11.19)$$

# Eigenvalues and eigenvectors

- Definition:
  - For an  $n \times n$  matrix  $A$ , an eigenvector  $v$  of  $A$  is a **nonzero** vector such that there exists some  $\lambda \in R$  satisfying
  - $Av = \lambda v$ .
- Discussion:
  - What's the dimension of  $v$ ?
  - Is  $Av$  a matrix or a vector?
  - Is  $\lambda v$  a vector or a real number?

# Goal: Find eigenvectors and eigenvalues of a square matrix $A$

- Discussion: How to find eigenvalues by hand?
  - Reviewed in Lecture 3.
- Work with the characteristic equation for the eigenvalues  $\lambda$ :

$$Av = \lambda v \Leftrightarrow Av - \lambda v = 0 \Leftrightarrow Av - \lambda I v = 0 \Leftrightarrow (A - \lambda I)v = 0.$$

- But  $v$  is a nonzero vector, so  $\det(A - \lambda I) = 0$ .

# Power method: Computing the eigenvalue of largest modulus and its corresponding eigenvector

$$|\lambda_1| > |\lambda_2| \geq \dots \geq |\lambda_n|.$$

- Works for diagonalizable matrix only. All symmetric matrices are diagonalizable.
- Algorithm:
  - Start with an initial nonzero vector  $w^{(0)}$
  - Run in K iterations

$$w^{(k+1)} = \frac{Aw^{(k)}}{\|Aw^{(k)}\|_2}.$$

- Then your final  $w^{(K)} \approx v_1$
- And  $\lambda_1 = Aw_1/v_1$

# Summary

- Power method is used to calculate eigenvalue and eigenvector of a matrix
  - In an iterative way
    - Stopping criterion: number of iterations, relative error
  - Works for diagonalizable matrices only
    - All symmetric matrices are diagonalizable
  - Only finds the eigenvalue of the largest absolute value and its associated eigenvector
    - HW2 also requires you to find the second largest eigenvalue. How?

# An example of D, L, U decomposition

- Linear system:

$$\begin{cases} 4x_1 + x_2 + 2x_3 = 4, \\ x_1 + 3x_2 + x_3 = 5, \\ 2x_1 + x_2 + 5x_3 = 6. \end{cases}$$

- Then,

$$A = \begin{bmatrix} 4 & 1 & 2 \\ 1 & 3 & 1 \\ 2 & 1 & 5 \end{bmatrix}, \quad x = \begin{bmatrix} x_1 \\ x_2 \\ x_3 \end{bmatrix}, \quad b = \begin{bmatrix} 4 \\ 5 \\ 6 \end{bmatrix}$$

$$D = \begin{bmatrix} 4 & 0 & 0 \\ 0 & 3 & 0 \\ 0 & 0 & 5 \end{bmatrix}, \quad L = \begin{bmatrix} 0 & 0 & 0 \\ 1 & 0 & 0 \\ 2 & 1 & 0 \end{bmatrix}, \quad U = \begin{bmatrix} 0 & 1 & 2 \\ 0 & 0 & 1 \\ 0 & 0 & 0 \end{bmatrix}$$

# Jacobi method

- After DLU decomposition, we have
  - $(D + L + U)x = b$
- Rearranging gives:
  - $Dx = b - (L + U)x$
- Jacobi iteration updates:
  - $x^{(k+1)} = D^{-1}(b - (L + U)x^{(k)})$ .
  - Or equivalently:
    - $x^{(k+1)} = D^{-1}b + \underbrace{(-D^{-1}(L + U))}_{=: T_J} x^{(k)}$ .
- So in compact form:
  - $x^{(k+1)} = T_J x^{(k)} + c$ , where  $c = D^{-1}b$ .

# Gauss-Seidel method

- After DLU decomposition, we have
  - $(D + L + U)x = b$
- Rearranging gives:
  - $(D + L)x = b - Ux$
- Gauss-Seidel iteration updates:
  - $(D + L)x^{(k+1)} = b - Ux^{(k)}$
  - Formally,  $x^{(k+1)} = T_{GS}x^{(k)} + c_{GS}$ 
    - where  $T_{GS} = -(D + L)^{-1}U, c_{GS} = (D + L)^{-1}b$

# Summary of Jacobi & Gauss-Seidel method

- Both use the L, U, D decomposition
  - $A = D + L + U$ ,
  - $D$ : diagonal of  $A$
  - $L$ : strictly lower triangular part of  $A$
  - $U$ : strictly upper triangular part of  $A$
- Both are guaranteed to converge if  $A$  is symmetric positive definite (SPD).
  - SPD:  $A = A^T$  and  $x^T A x > 0$  if any  $x \neq 0$ .

# Matrix rank of $A$ ( $m \times n$ matrix)

- Definition:
  - Maximal number of **linearly independent columns** or maximal number of linearly independent **rows**.
- Two examples: Find ranks using Gaussian Elimination.

$$A = \begin{bmatrix} 1 & 0 & 0 \\ 0 & 2 & 0 \\ 0 & 0 & 3 \end{bmatrix}$$

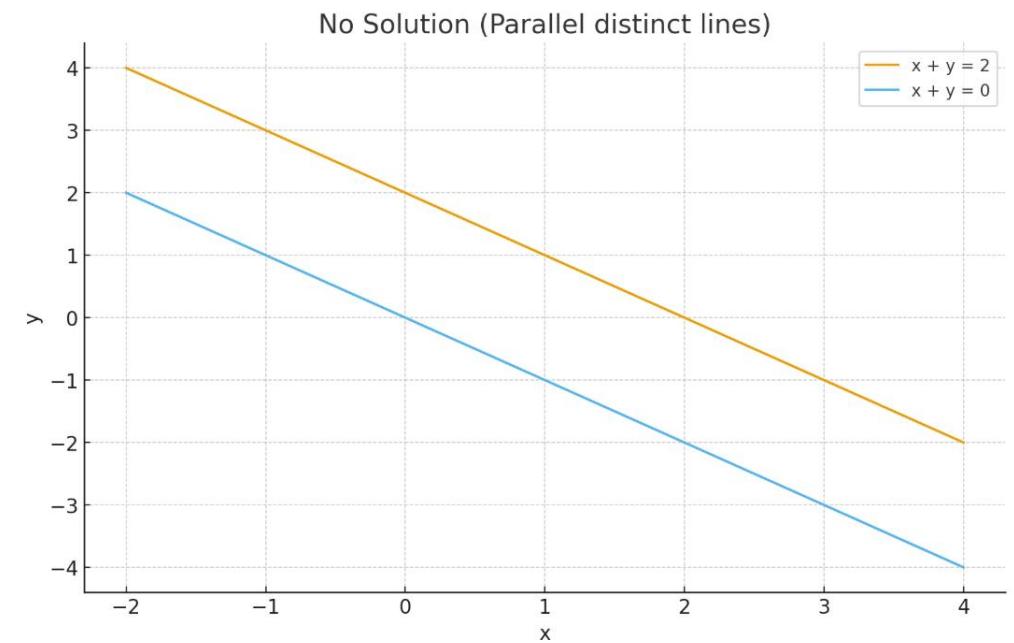
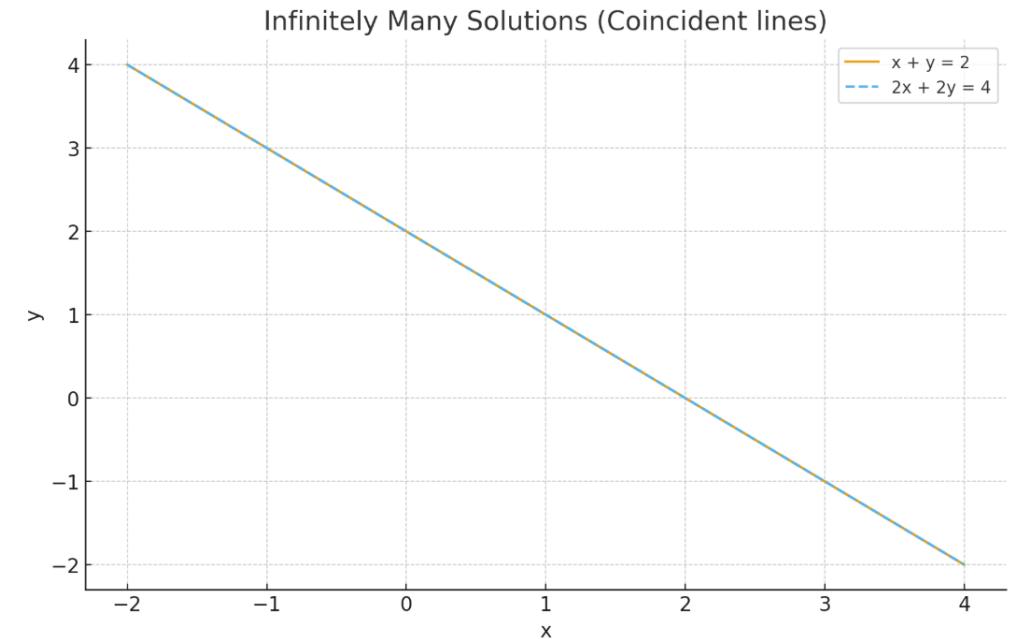
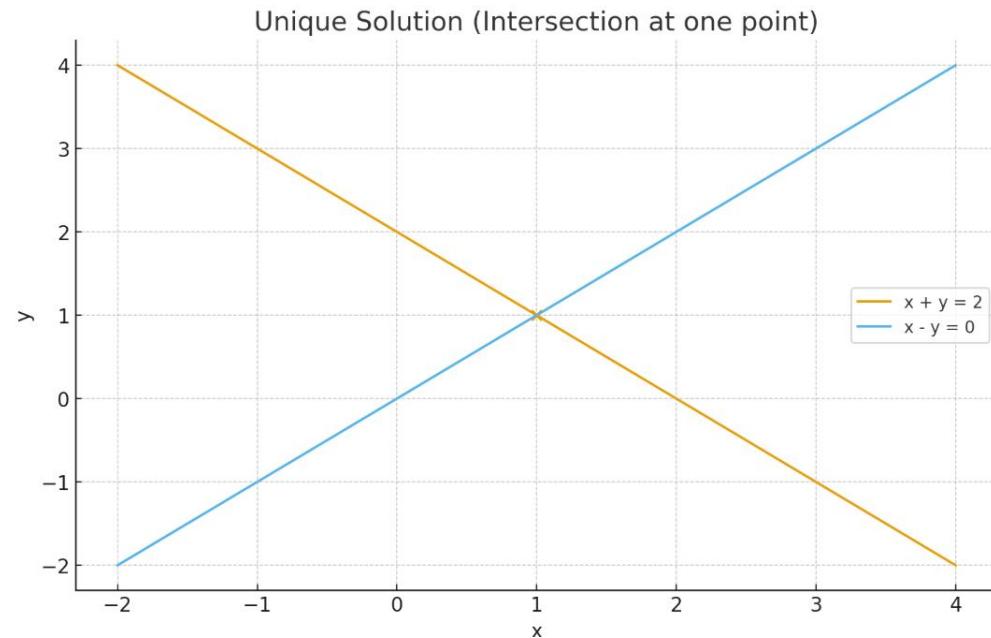
$$B = \begin{bmatrix} 1 & 2 & 3 \\ 2 & 4 & 6 \\ 0 & 1 & 1 \end{bmatrix} \xrightarrow{R_2 \leftarrow R_2 - 2R_1} \begin{bmatrix} 1 & 2 & 3 \\ 0 & 0 & 0 \\ 0 & 1 & 1 \end{bmatrix}$$

- Rank is number of pivots after Gaussian Elimination.

# Summary of Conditions for Solutions of a Linear System $Ax = b$

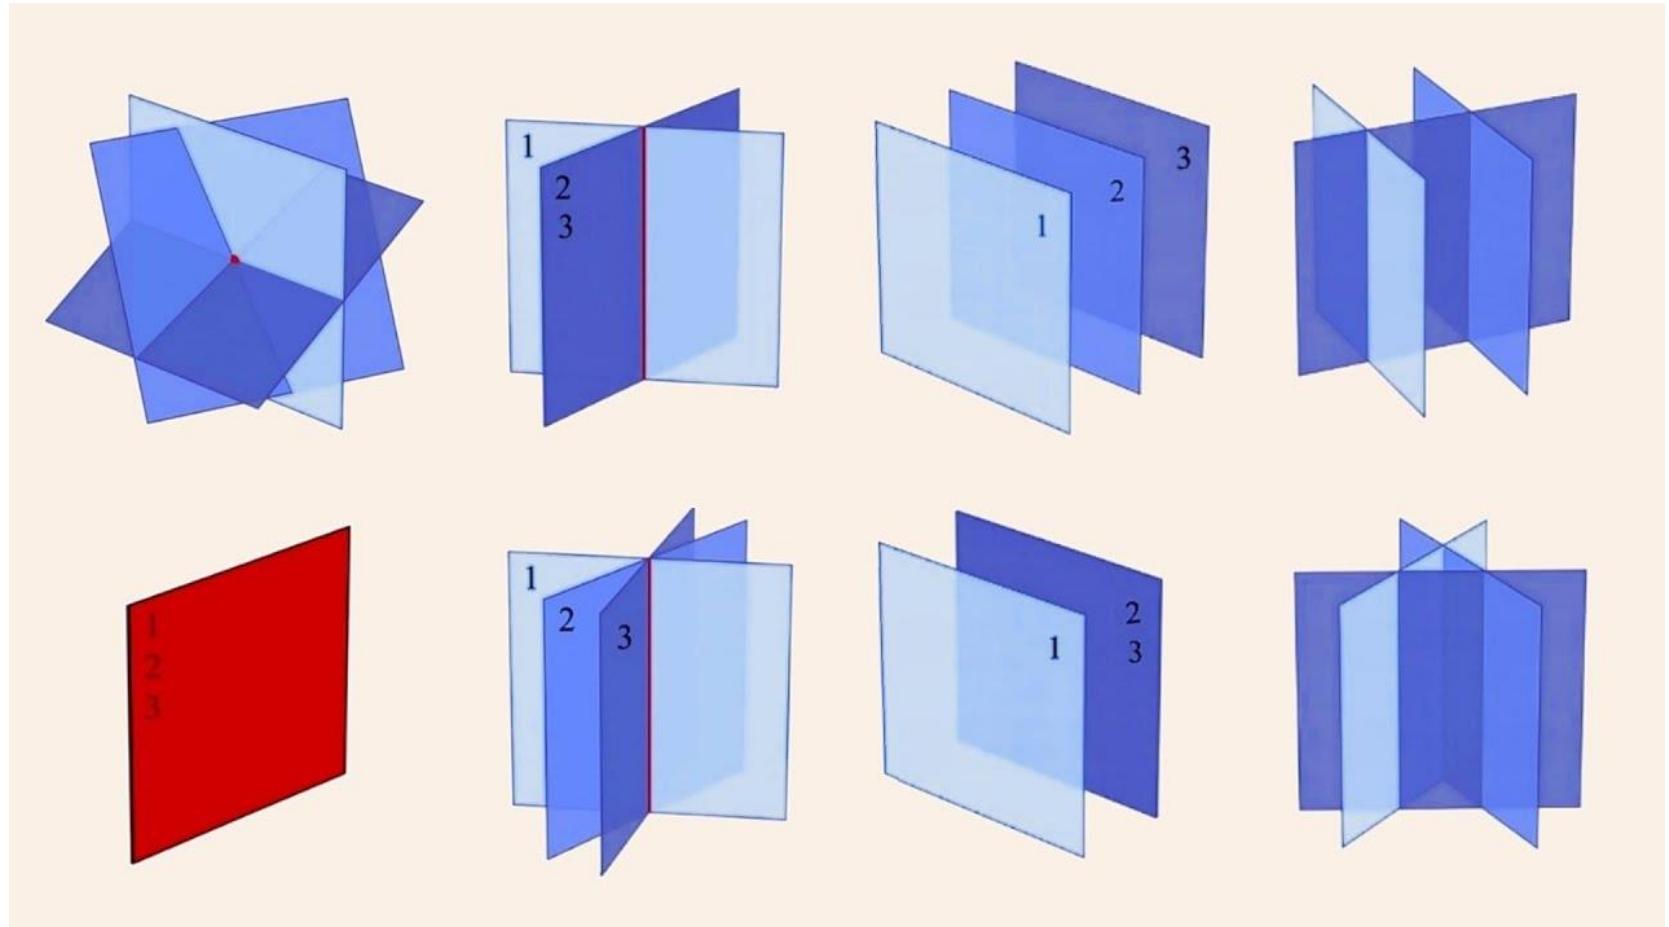
Case	Rank Condition	Number of Solutions	Geometric Interpretation
No Solution	$\text{rank}(A) < \text{rank}([A   \mathbf{b}])$	None (inconsistent system)	Hyperplanes do not intersect (contradictory equations)
Unique Solution	$\text{rank}(A) = \text{rank}([A   \mathbf{b}]) = n$	Exactly one	Hyperplanes intersect at a single point
Infinitely Many Solutions	$\text{rank}(A) = \text{rank}([A   \mathbf{b}]) < n$	Infinitely many	Hyperplanes intersect along a line, plane, or higher-dimensional subspace

# Geometric view of these three systems



# Geometric view of solutions (3-d case)

- Discussion:
  - Which figure shows a *unique* solution?
  - Which figure shows *infinitely many* solutions?
  - Which figure shows *no* solution?



# Case study: Housing price

- Suppose we would like to build a model predicting house prices.
  - The model takes **features of a house** as inputs, and outputs **predicted price**.
- Discussion:
  - What are the factors (features) of a house that affects its price?
- For example,
  - 8 features:

- MedInc	median income in block group
- HouseAge	median house age in block group
- AveRooms	average number of rooms per household
- AveBedrms	average number of bedrooms per household
- Population	block group population
- AveOccup	average number of household members
- Latitude	block group latitude
- Longitude	block group longitude
  - 1 label: house price

# Linear model

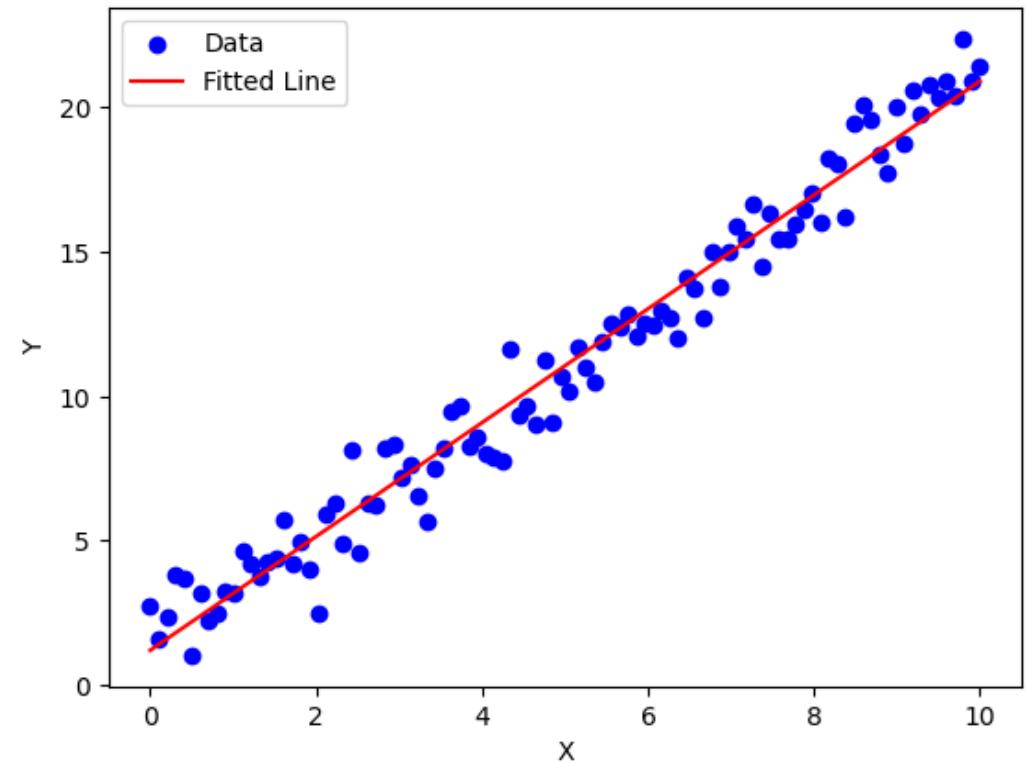
- Take input feature vector
  - $\text{Price}(x) = w_1x_1 + w_2x_2 + w_3x_3 + w_4x_4 + \dots$
  - $x_1$ : median income
  - $x_2$ : median house age
  - $x_3$ : average number of rooms
  - $x_4$ : average number of bedrooms
  - ...
- Label space is the real number space  $R$

# Linear model

- In vector form:
  - $\text{Price}(x) = x^T w$
  - $x = [x_1, x_2, \dots, x_8]$ : feature vector
  - $w = [w_1, w_2, \dots, w_8]$ : parameter vector
- As long as we find a good  $w$ , we have a good linear model.
- Goal: Find a good  $w$ .

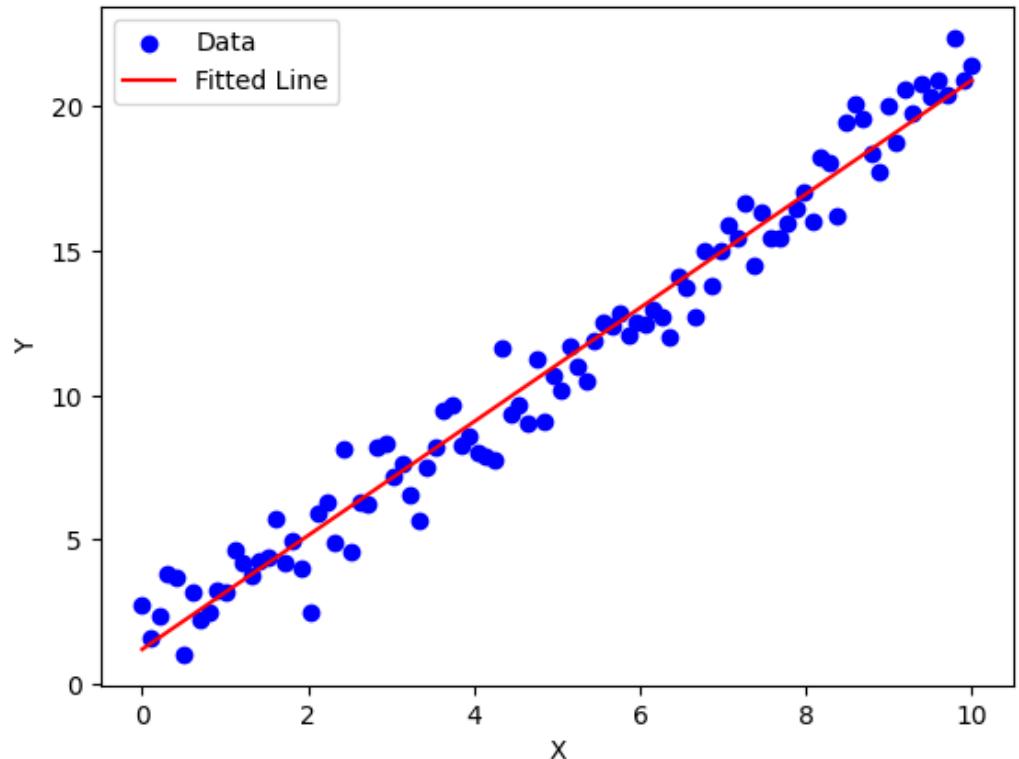
# Considering conditions of linear systems

- In real-world applications, there are many **challenges**.
  - No solution
  - Noisy data
  - Overdetermined systems (most common case)
    - Fitting a hyperplane (a line in 2-d) to too many data points.
- Right figure:
  - $x$  is a feature of the house
  - $y$  is the price.



# Considering conditions of linear systems

- In real-world applications, there are many **challenges**.
  - No solution
  - Noisy data
  - Overdetermined systems (most common case)
    - Fitting a hyperplane (a line in 2-d) to too many data points.
- So our goal reduces to find the an approximate  $w$  that **best describes** the data!
- How?



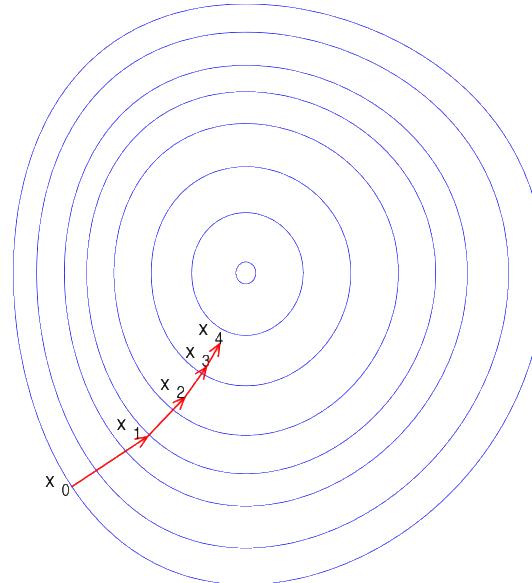
# The objective function for learning linear regression under **square loss**

- $\hat{w} = \operatorname{argmin}_w \frac{1}{n} \sum_{i=1}^n (x_i^T w - y_i)^2 = \operatorname{argmin}_w \|Xw - y\|_2^2$ 
  - aka: Ordinary Least Square (OLS)
  - In-class exercise: solve this optimization problem by setting gradient of the objective function to 0.

# How do we optimize a continuously differentiable function in general?

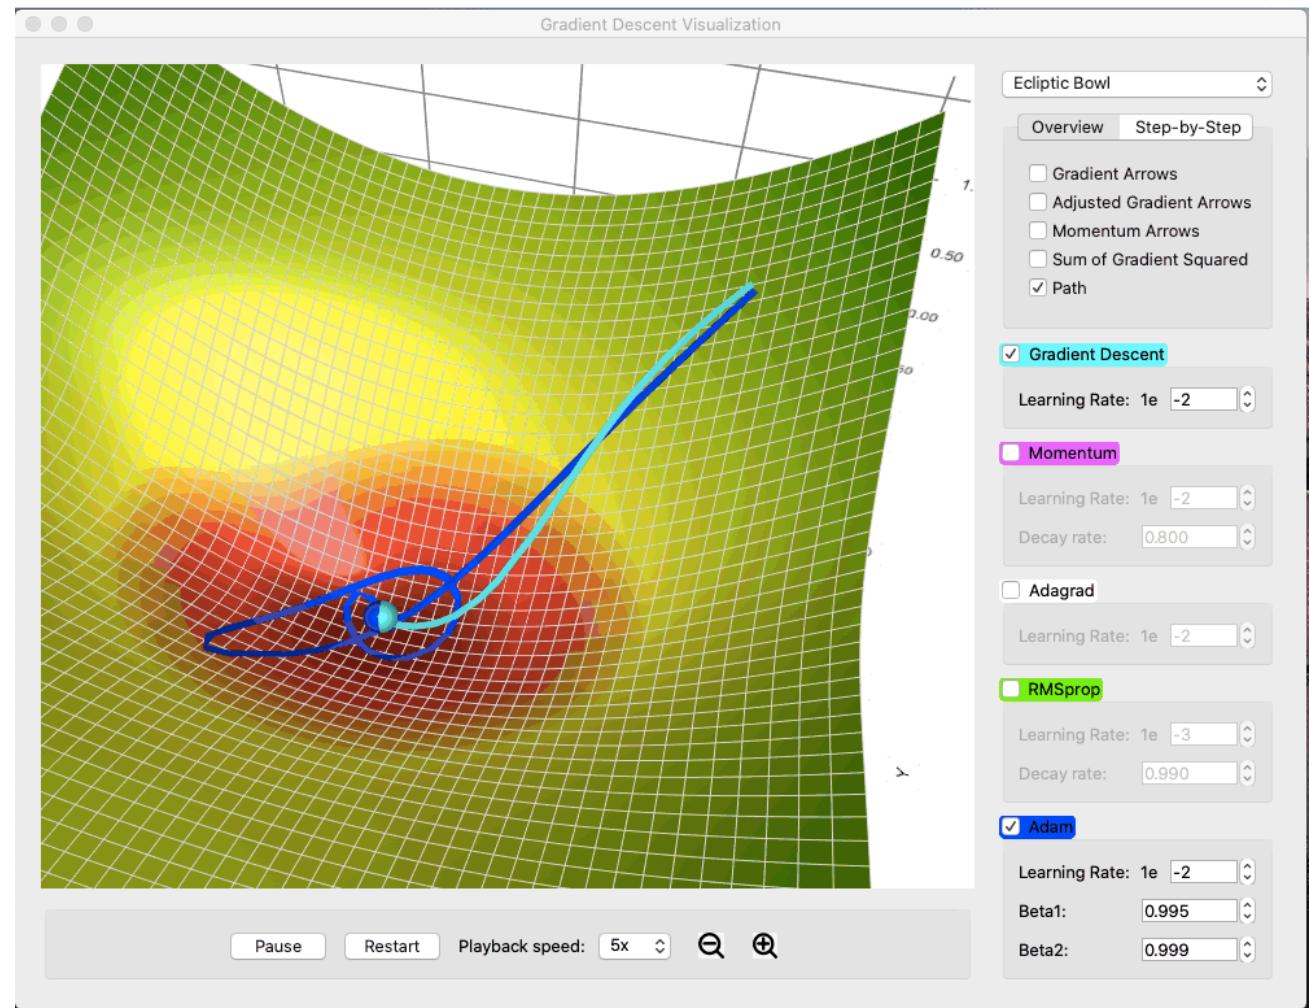
- The problem:  $\min_{\theta} f(\theta)$
- Discussion: How do you solve this optimization problem?
- Gradient descent in iterations

$$\theta_{t+1} = \theta_t - \eta_t \nabla f(\theta_t)$$



# Gradient Descent Demo in 2-D

- An excellent demo tool:
  - [https://github.com/lilipads/gradient\\_descent\\_viz](https://github.com/lilipads/gradient_descent_viz)



# Gradient descent for quadratic function

- $\min f(x) = x^2$
- Follow me on the first two examples:
  1. Find  $x_4$  given  $x_0 = 2, \eta = 0.1$
  2. Find  $x_4$  given  $x_0 = 2, \eta = 0.4$
- In-class exercise questions:
  1. Find  $x_4$  given  $x_0 = 4, \eta = 0.4$
  2. Find  $x_4$  given  $x_0 = 2, \eta = 1.5$  (what did you find?)

# Back to linear regression: How to solve it using Gradient Descent?

- $\hat{w} = \operatorname{argmin}_w \frac{1}{n} \sum_{i=1}^n (x_i^T w - y_i)^2 = \operatorname{argmin}_w \|Xw - y\|_2^2$
- In-class exercise: Write the GD updating rule for solving  $w$ .
  - $w \leftarrow w - 2\eta X^T(Xw - y)$

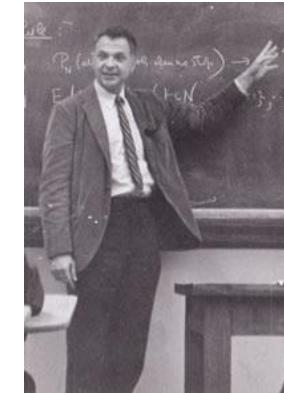
# Checkpoint

- Least square:
  - Heavily used in practice, due to
    - Large datasets (many data points)
    - Noisy data
    - No solution based on conditions of linear systems
- Linear regression
  - $\hat{w} = \operatorname{argmin}_w \frac{1}{n} \sum_{i=1}^n (x_i^T w - y_i)^2 = \operatorname{argmin}_w \|Xw - y\|_2^2$
  - Direct solver:  $\hat{w} = (X^T X)^{-1} X^T y$
  - GD:  $w \leftarrow w - 2\eta X^T (Xw - y)$

# Stochastic Gradient Descent (Robbins-Monro 1951)

- Gradient descent

$$\theta_{t+1} = \theta_t - \eta_t \nabla f(\theta_t)$$



Herbert Robbins  
1915 - 2001

- Stochastic gradient descent
  - Using a **stochastic approximation** of the gradient:

$$\theta_{t+1} = \theta_t - \eta_t \hat{\nabla} f(\theta_t)$$

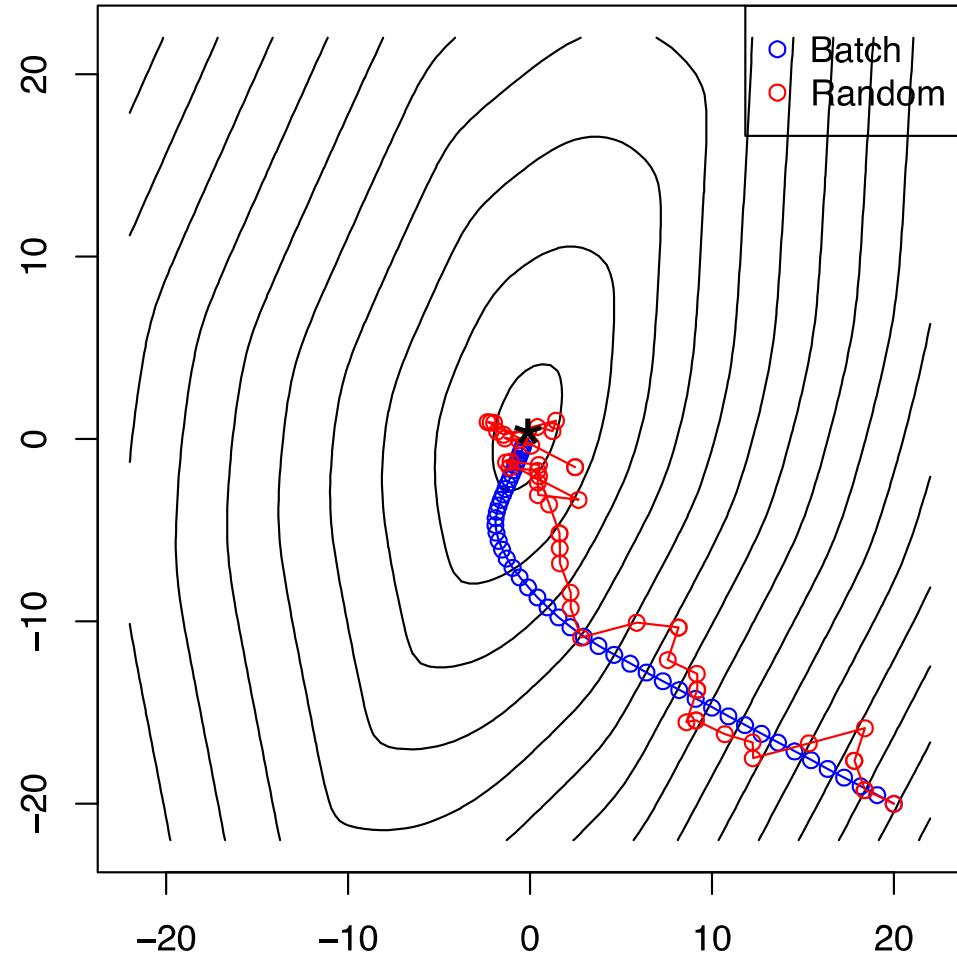
# A natural choice of SGD in machine learning

- Recall that

$$\min_{\theta \in \mathbb{R}^d} \frac{1}{n} \sum_{i=1}^n \ell(\theta, (x_i, y_i))$$

- SGD samples a data point  $i$  uniformly at random while GD uses all data!
  - Use  $\nabla_{\theta} \ell(\theta, (x_i, y_i))$

# Illustration of **GD** vs **SGD**



Time complexity:

GD:  $O(nd * n_{iterations})$   
SGD:  $O(d * n_{iterations})$

# The power of SGD

- Extremely simple:
  - A few lines of code
- Extremely scalable
  - Just a few pass of the data, no need to store the data
- Extremely general:
  - In addition to linear regression, in practice it can solve most optimization problems of differentiable functions
    - E.g., Training neural networks, Transformer, Generative Pretrained Transformer
  - **Foundational** algorithm of the AI revolution as we see today!

# Time complexity of direct solver and GD/SGD for solving linear regression

- Direct solver
  - $O(nd^2 + d^3)$
- GD:
  - $O(ndT)$
- SGD:
  - $O(dT)$
- $T = \text{n\_iterations}$

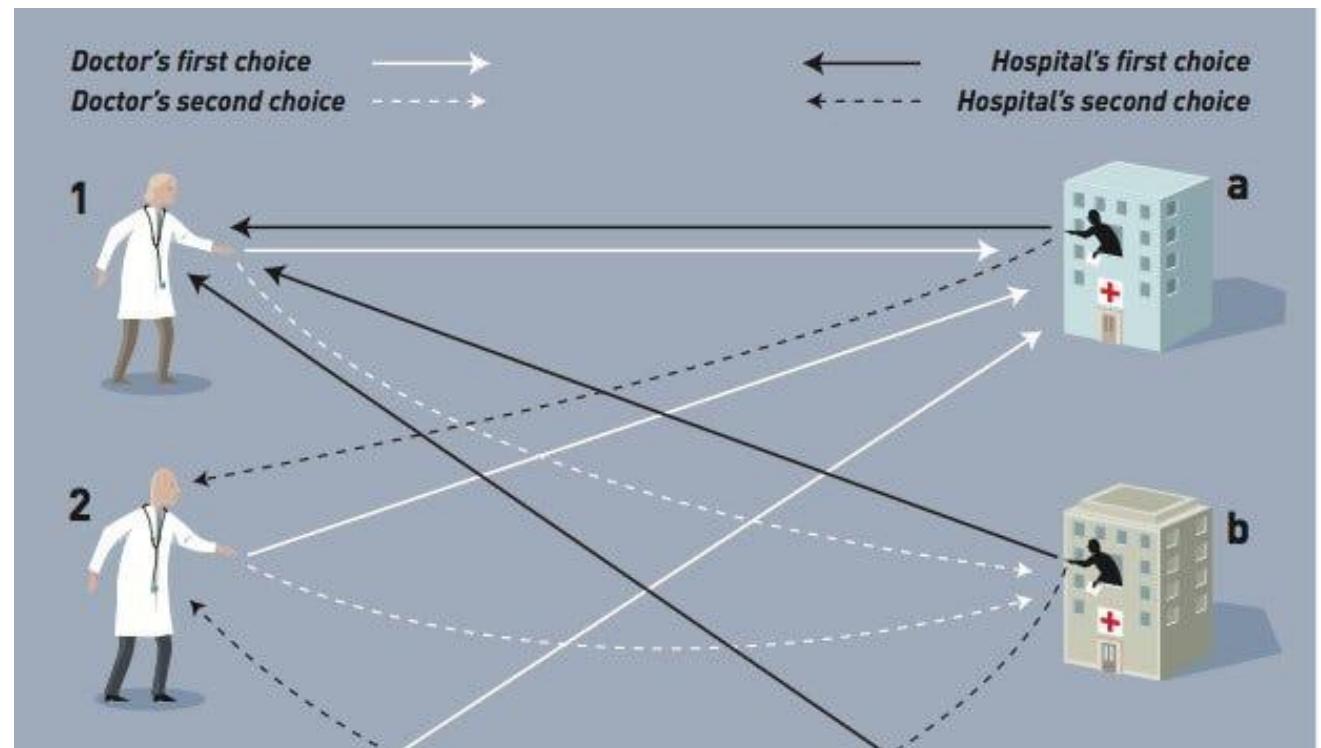
# What's Linear Programming (LP)?

- An optimization problem of **linear** objective functions with **linear** constraints.
  - Objective function can be minimized or maximized
  - Constraints can be in equalities or inequalities
  - All functions must be linear functions
- 2 examples:

$\min x_1 + 2x_2$	$\max x_1 + 2x_2$
s.t. $x_1 + x_2 \leq 3$	s.t. $x_1 + x_2 = 3$
$x_1 \geq 1$	
- Discussion: Could you propose more linear programming problems?

# Application of LP: Matching problem

- Company (hospital) - Candidate (doctor) matching problem
- Each doctor:
  - Fits one position
- Doctors/hospitals:
  - Have their preferences
- Goal:
  - Put doctors to positions
  - Such that overall best match



# Application of LP: Optimal transport

- Suppose you run a company, which has 4 factories and 3 big markets, each in a different city.

Buffalo supplies 15	Syracuse supplies 15	Boston demands 25
Rochester supplies 20	Albany supplies 35	
New York City demands 30		
Philadelphia demands 30		

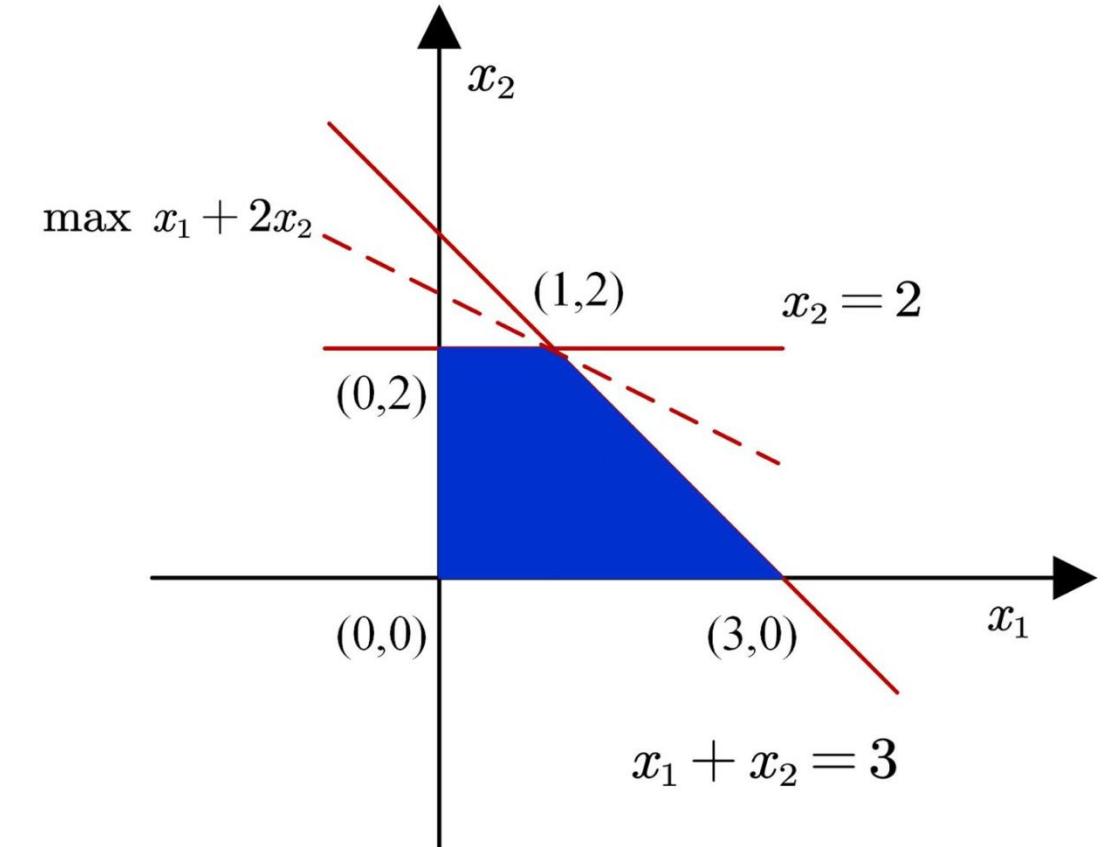
- Your job is to design the optimal transportation route that has minimum transportation cost of your products
  - Each route (supply to demand) costs differently
  - Each factory has its supply capacity
  - Each market must be well supplied to maximize your profit

# How to solve the LP problem?

maximize  $x_1 + 2x_2$   
subject to  $x_1 + x_2 \leq 3$

$$\begin{aligned}x_2 &\leq 2 \\x_1 &\geq 0 \\x_2 &\geq 0\end{aligned}$$

- For most 2-d LP problems,
  1. We can draw it's feasible region
  2. And move it's objective function
- In-class exercise: Draw the feasible region defined by constraints.



# From primal to dual LP problems

- In-class exercise:
  - Work on the following two LP problems by drawing graphs

$$\max Z = x_1 + x_2$$

$$\text{s.t. } \begin{cases} x_1 + 2x_2 \leq 8, \\ 3x_1 + 2x_2 \leq 18, \\ x_1, x_2 \geq 0. \end{cases}$$

$$\min Z = 8y_1 + 18y_2$$

$$\text{s.t. } \begin{cases} y_1 + 3y_2 \geq 1, \\ 2y_1 + 2y_2 \geq 1, \\ y_1, y_2 \geq 0. \end{cases}$$

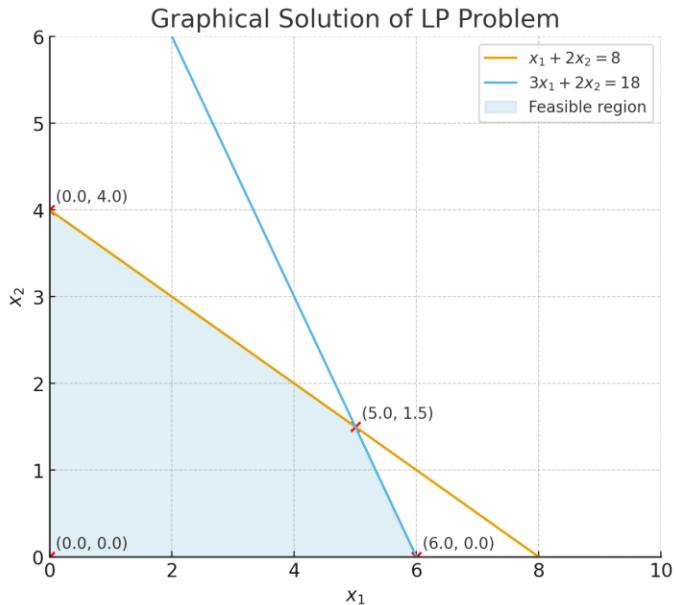
- What can you see from their optimal  $Z$ ?

# From primal to dual LP problems

- Solutions to in-class exercise:

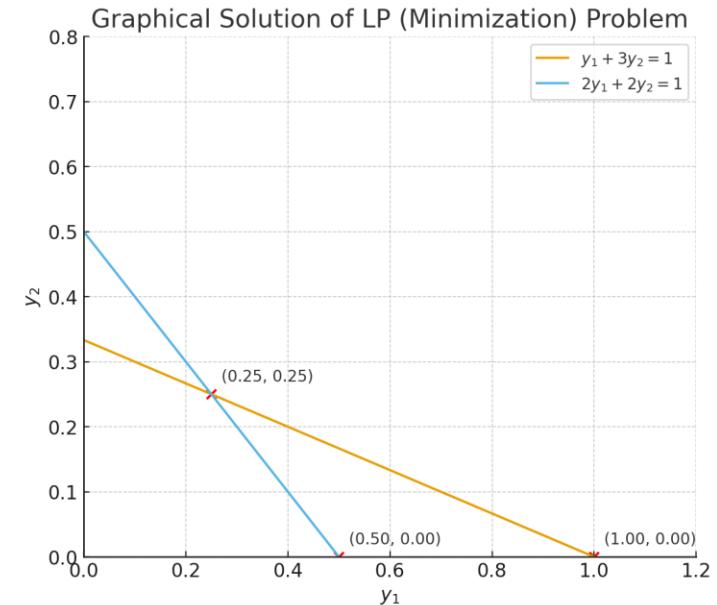
$$\max Z = x_1 + x_2$$

$$\text{s.t. } \begin{cases} x_1 + 2x_2 \leq 8, \\ 3x_1 + 2x_2 \leq 18, \\ x_1, x_2 \geq 0. \end{cases}$$



$$\min Z = 8y_1 + 18y_2$$

$$\text{s.t. } \begin{cases} y_1 + 3y_2 \geq 1, \\ 2y_1 + 2y_2 \geq 1, \\ y_1, y_2 \geq 0. \end{cases}$$



- They are primal and dual LP problems!

# From primal to dual LP problems

- Primal problem:

$$\max z = 4x_1 + x_2 + 5x_3 + 3x_4$$

$$\begin{aligned}x_1 - x_2 - x_3 + 3x_4 &\leq 1 \\5x_1 + x_2 + 3x_3 + 8x_4 &\leq 55 \\-x_1 + 2x_2 + 3x_3 - 5x_4 &\leq 3 \\x_i &\geq 0\end{aligned}$$

- Key idea:

- Multiply each constraint with a non-negative multiplier and form linear combinations of constraints.

$$y_1(x_1 - x_2 - x_3 + 3x_4) + y_2(5x_1 + x_2 + 3x_3 + 8x_4) + y_3(-x_1 + 2x_2 + 3x_3 - 5x_4) \leq y_1 + 55y_2 + 3y_3.$$

$$(y_1 + 5y_2 - y_3)x_1 + (-y_1 + y_2 + 2y_3)x_2 + (-y_1 + 3y_2 + 3y_3)x_3 + (3y_1 + 8y_2 - 5y_3)x_4 \leq y_1 + 55y_2 + 3y_3.$$

- Finally, dual problem:

$$\begin{aligned}\min u &= y_1 + 55y_2 + 3y_3 & y_1 + 5y_2 - y_3 &\geq 4 \\&& -y_1 + y_2 + 2y_3 &\geq 1 \\&& -y_1 + 3y_2 + 3y_3 &\geq 5 \\&& 3y_1 + 8y_2 - 5y_3 &\geq 3 \\&& y_i &\geq 0\end{aligned}$$

# Dual problem of linear programming

- Economic Interpretation
  - The dual variables  $y$  represent **shadow prices** — the value of relaxing each constraint by one unit.
  - In a resource allocation problem, each  $y_i$  tells how much the objective (profit) would improve if resource  $i$  were increased slightly.
- Weak Duality:  
For any feasible  $x$  (primal) and  $y$  (dual),  $c^T x \leq b^T y$ .
  - The dual provides an **upper bound** (for maximization problems).
- Strong Duality:  
At the optimal solutions  $x^*, y^*$ ,  $c^T x^* = b^T y^*$ .
  - Solving one problem solves the other — they share the **same** optimal value.

# Dual problem of linear programming

- Why we study dual problems?
- Duality helps:
  - **Check optimality:** If primal and dual feasible solutions give the same objective, both are optimal.
  - **Perform sensitivity analysis:** Dual variables show how changes in constraints affect the outcome.
  - **Simplify computation:** Some LPs are easier to solve in dual form (e.g., when constraints >> variables).

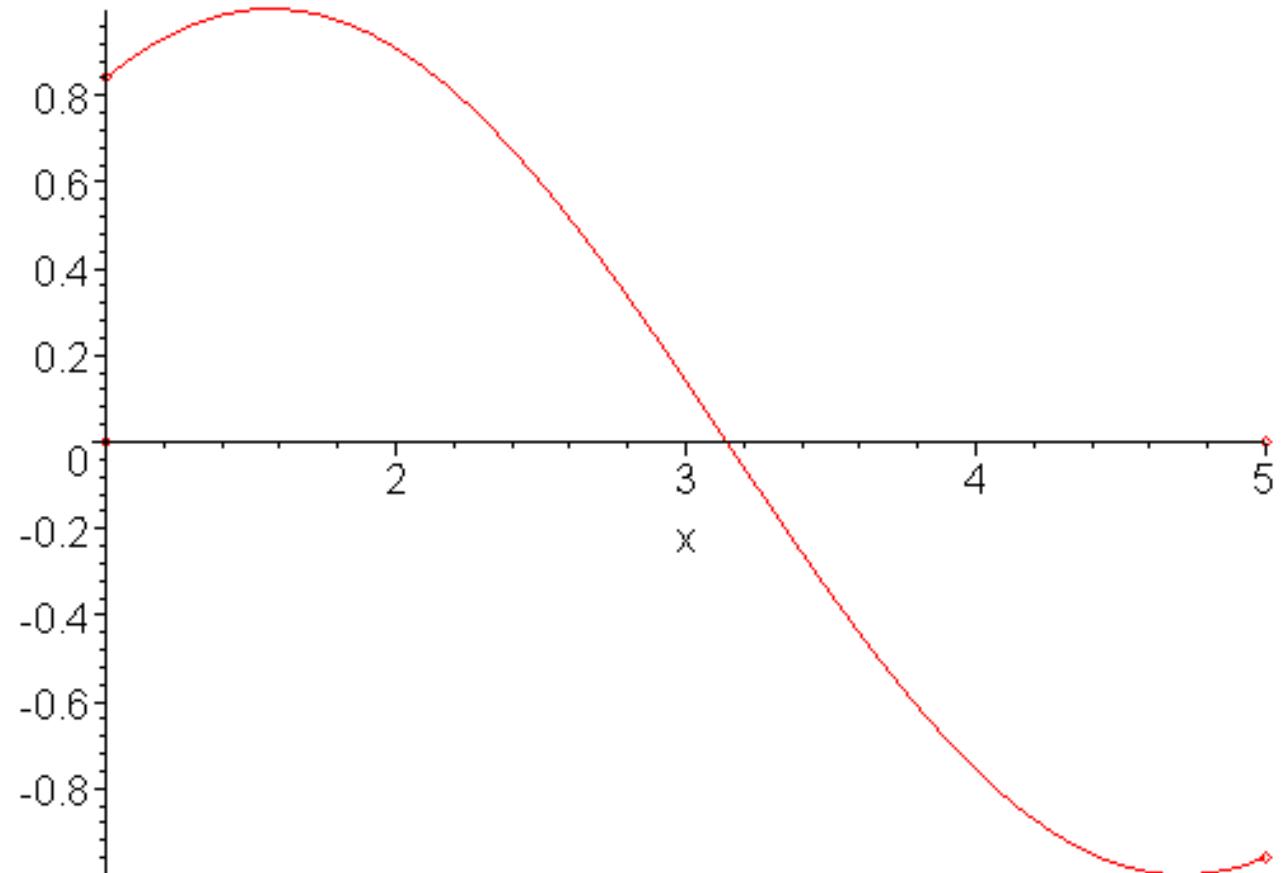
# Nonlinear equation solver: Bisection method

- Key idea:
  - In every iteration, we cut the interval in half while still maintaining the property that the **endpoints have opposite signs**. This allows us to conclude that we're getting closer and closer to a root.
- Algorithm:
  1. Preprocessing: If  $F(a) = 0$  or  $F(b) = 0$ , output whichever one was 0 and terminate. If  $F(a) < 0 < F(b)$ , then set  $inc = 1$ . Otherwise, set  $inc = 0$ .
  2. Compute  $z = \frac{a+b}{2}$ , the midpoint of the interval  $[a, b]$ .
  3. If  $F(z) = 0$ , return  $z$  and terminate.
  4. If  $inc = 0$  (so  $F(a) > 0 > F(b)$ ):
    - (a) If  $F(z) < 0$ , then set  $b = z$ .
    - (b) If  $F(z) > 0$ , then set  $a = z$ .
  5. If  $inc = 1$  (so  $F(a) < 0 < F(b)$ ):
    - (a) If  $F(z) < 0$ , then set  $b = z$ .
    - (b) If  $F(z) > 0$ , then set  $a = z$ .

After  $k$  iterations, we output the midpoint of the resulting interval.

# Illustration of the bisection method

- Initial interval: [1, 5]
- 3 steps in each iteration:
  - Given  $a, b$ , find midpoint
  - Check midpoint value
  - Update  $a$  or  $b$



# Nonlinear equation solver: Newton's method

- Key idea:
  - Take  $F$ , find its local linear approximation at a starting point  $x_0$ , solve for  $x$  to get  $x_1$ , and use that as our new initial point.
  - Iterate until (hopefully) convergence.
- So, how to find the local linear approximation of  $F$  at  $x_0$ ?
  - First-order Taylor expansion at  $x_0$

$$P_1(x) = F(x_0) + F'(x_0)(x - x_0).$$

$$0 = F(x_0) + F'(x_0)(x - x_0) \implies -\frac{F(x_0)}{F'(x_0)} = x - x_0 \implies x = x_0 - \frac{F(x_0)}{F'(x_0)}.$$

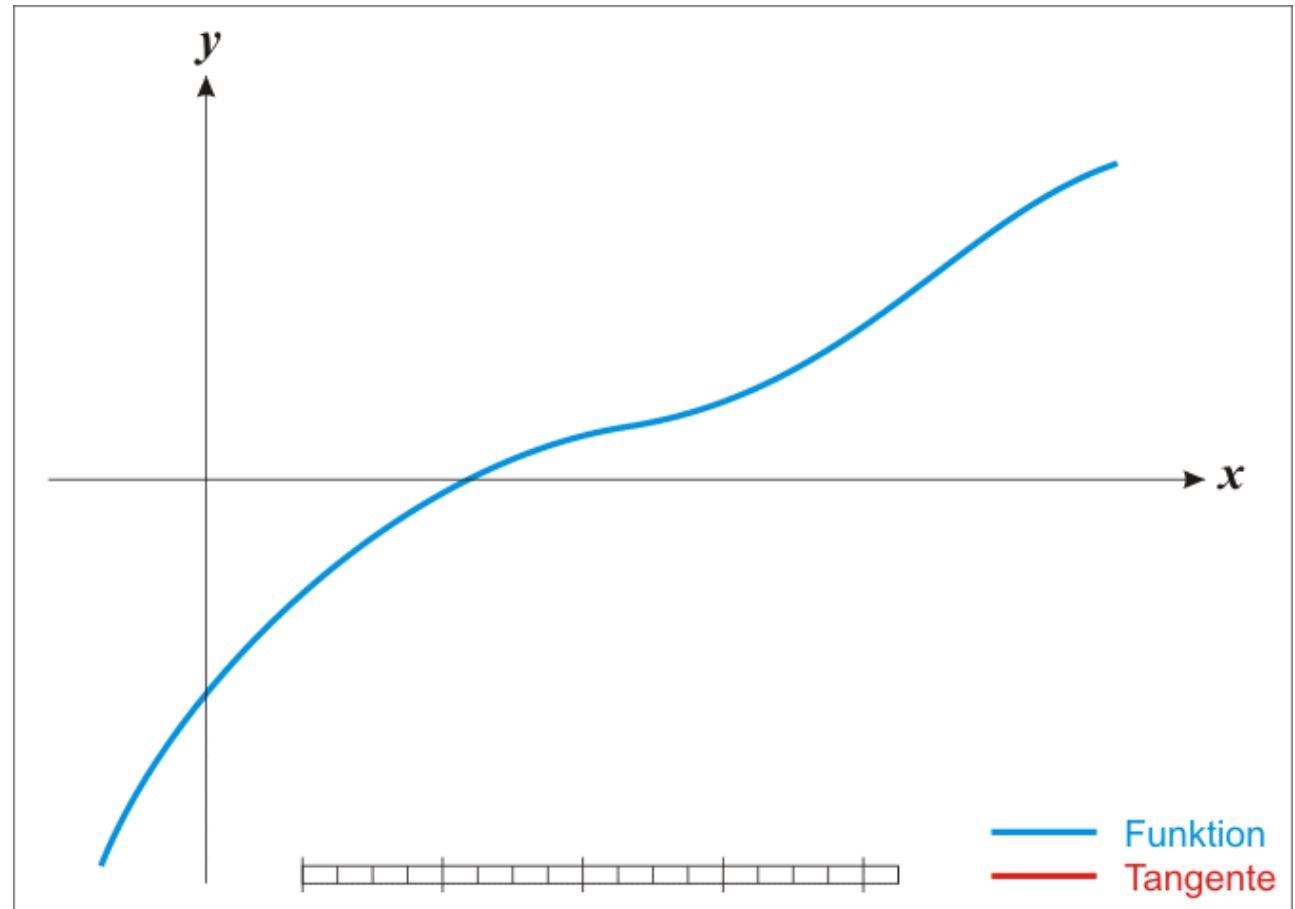
- Algorithm: (Newton update equation)

$$x_{k+1} = x_k - \frac{F(x_k)}{F'(x_k)}.$$

# Illustration of the Newton's method

- In each iteration:

$$x_{k+1} = x_k - \frac{F(x_k)}{F'(x_k)}.$$



# Summary of nonlinear equation solvers

- Things to know about:
  - Problem statement
  - Assumptions behind each method
  - Benefits/drawbacks of each method
  - Key theorems from calculus that feature in their analysis
  - How does each method look, visually?
  - How do we code each method up in Matlab/Python?
- Technical summary table:

Methods	Bisection method	Newton's method
Assumptions	Continuity, opposite sign condition	Continuous, differentiable, initial point close to root
Associated theorem	Intermediate value theorem	Taylor's remainder theorem
Guarantee	Linear convergence	Quadratic convergence

# Problem setup of Interpolation

- For given data
  - $(t_1, y_1), (t_2, y_2), \dots, (t_m, y_m)$  with  $t_1 < t_2 < \dots < t_m$
- determine function  $f: R \rightarrow R$  such that
  - $f(t_i) = y_i, \forall i = 1, \dots, m$
  - Exactly crossing all data points!
- $f$  is **interpolating function**, or **interpolant**, for given data.
  - $f$  could be function of more than one variable, but let's focus on the 1-dimensional case first.

# Interpolation vs. Regression

- By definition, interpolating function fits given data points exactly
- Interpolation is inappropriate if data points subject to significant errors
  - Regression is a better choice in this case
- It is usually preferable to smooth noisy data
- Regression is more appropriate for special function libraries
  - Linear regression

# Basis Functions

- Family of functions for interpolating:
  - Set of basis functions  $\phi_1(t), \dots, \phi_n(t)$
- Interpolating function  $f$  is chosen as linear combination of them

$$f(t) = \sum_{j=1}^n x_j \phi_j(t)$$

- Requiring  $f$  to interpolate data  $(t_i, y_i)$  means

$$f(t_i) = \sum_{j=1}^n x_j \phi_j(t_i) = y_i, \quad i = 1, \dots, m$$

- Discussion: What is this system?
  - A system of linear equations  $Ax = y$  for  $n$ -vector  $x$  of parameters  $x_j$ , where entries of  $m \times n$  matrix  $A$  are given by  $a_{ij} = \phi_j(t_i)$ .

# Basic polynomial interpolation

- Simplest and most common type of interpolation using polynomials
- Unique polynomial of degree at most  $n - 1$  passes through  $n$  data points  $(t_i, y_i), i = 1, \dots, n$ , where  $t_i$  are distinct

# Basic polynomial interpolation

- Basis functions

$$\phi_j(t) = t^{j-1}, \quad j = 1, \dots, n$$

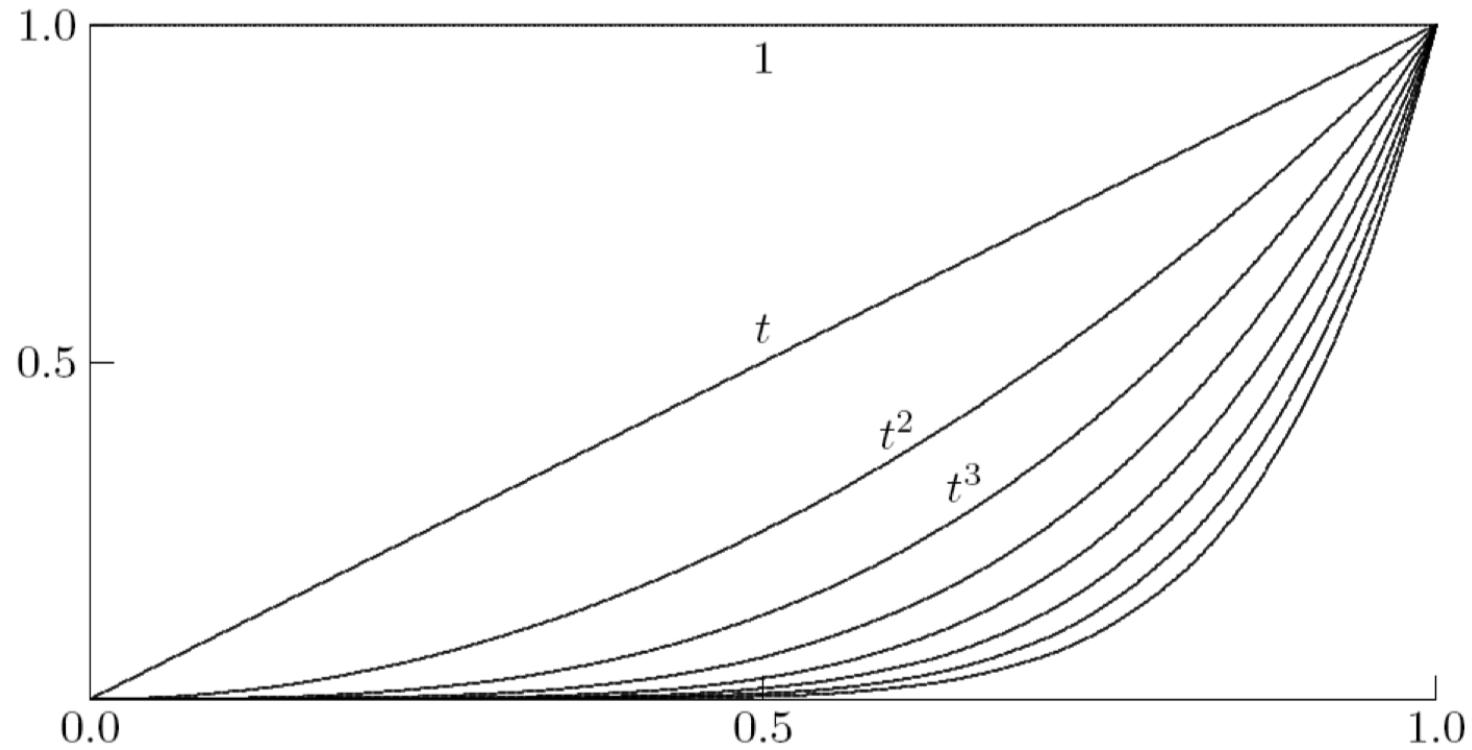
- give interpolating polynomial of form

$$p_{n-1}(t) = x_1 + x_2 t + \cdots + x_n t^{n-1}$$

- with coefficients  $x$  given by  $n \times n$  linear system

$$\mathbf{A}\mathbf{x} = \begin{bmatrix} 1 & t_1 & \cdots & t_1^{n-1} \\ 1 & t_2 & \cdots & t_2^{n-1} \\ \vdots & \vdots & \ddots & \vdots \\ 1 & t_n & \cdots & t_n^{n-1} \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_n \end{bmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ \vdots \\ y_n \end{bmatrix} = \mathbf{y}$$

# Basis functions



# Lagrange interpolation

- For given set of data points  $(t_i, y_i), i = 1, \dots, n$ , let

$$\ell(t) = \prod_{k=1}^n (t - t_k) = (t - t_1)(t - t_2) \cdots (t - t_n)$$

- Define weights

$$w_j = \frac{1}{\ell'(t_j)} = \frac{1}{\prod_{k=1, k \neq j}^n (t_j - t_k)}, \quad j = 1, \dots, n$$

- Lagrange basis functions are then given by

$$\ell_j(t) = \ell(t) \frac{w_j}{t - t_j}, \quad j = 1, \dots, n$$

- From definition,  $\ell_j(t)$  is polynomial of degree  $n - 1$

# Lagrange interpolation

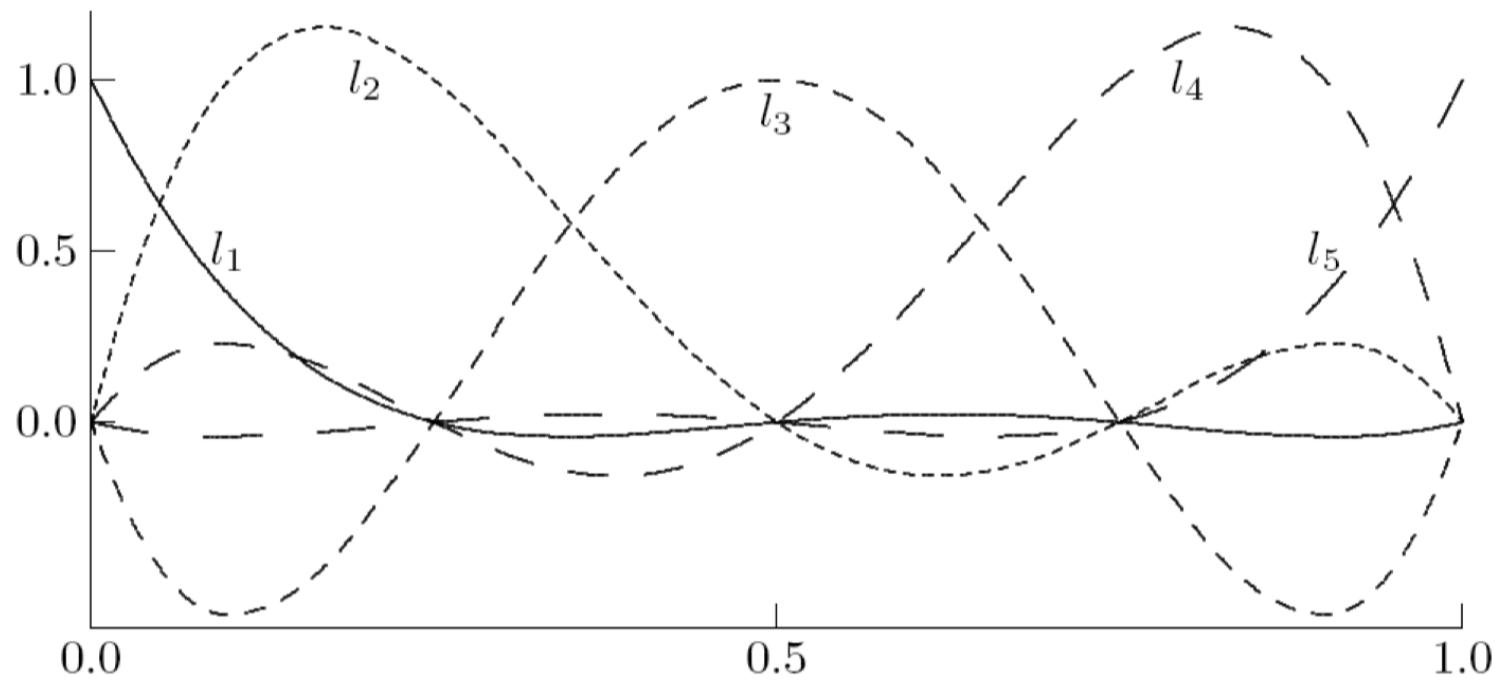
- Assuming common factor  $(t_i - t_j)$  in  $\ell(t_j)/(t_i - t_j)$  is canceled to avoid division by zero when evaluating  $\ell_j(t_i)$ , then

$$\ell_j(t_i) = \begin{cases} 1 & \text{if } i = j \\ 0 & \text{if } i \neq j \end{cases}, \quad i, j = 1, \dots, n$$

- Matrix of linear system  $Ax = y$  is identity matrix  $I$
- Coefficients  $x$  for Lagrange basis functions are just data values  $y$
- Polynomial of degree  $n - 1$  interpolating data points  $(t_i, y_i), i = 1, \dots, n$  is given by

$$p_{n-1}(t) = \sum_{j=1}^n y_j \ell_j(t) = \sum_{j=1}^n y_j \ell(t) \frac{w_j}{t - t_j} = \ell(t) \sum_{j=1}^n y_j \frac{w_j}{t - t_j}$$

# Lagrange Basis Functions



# Newton interpolation

- For given set of data points  $(t_i, y_i), i = 1, \dots, n$ , Newton basis functions are defined by

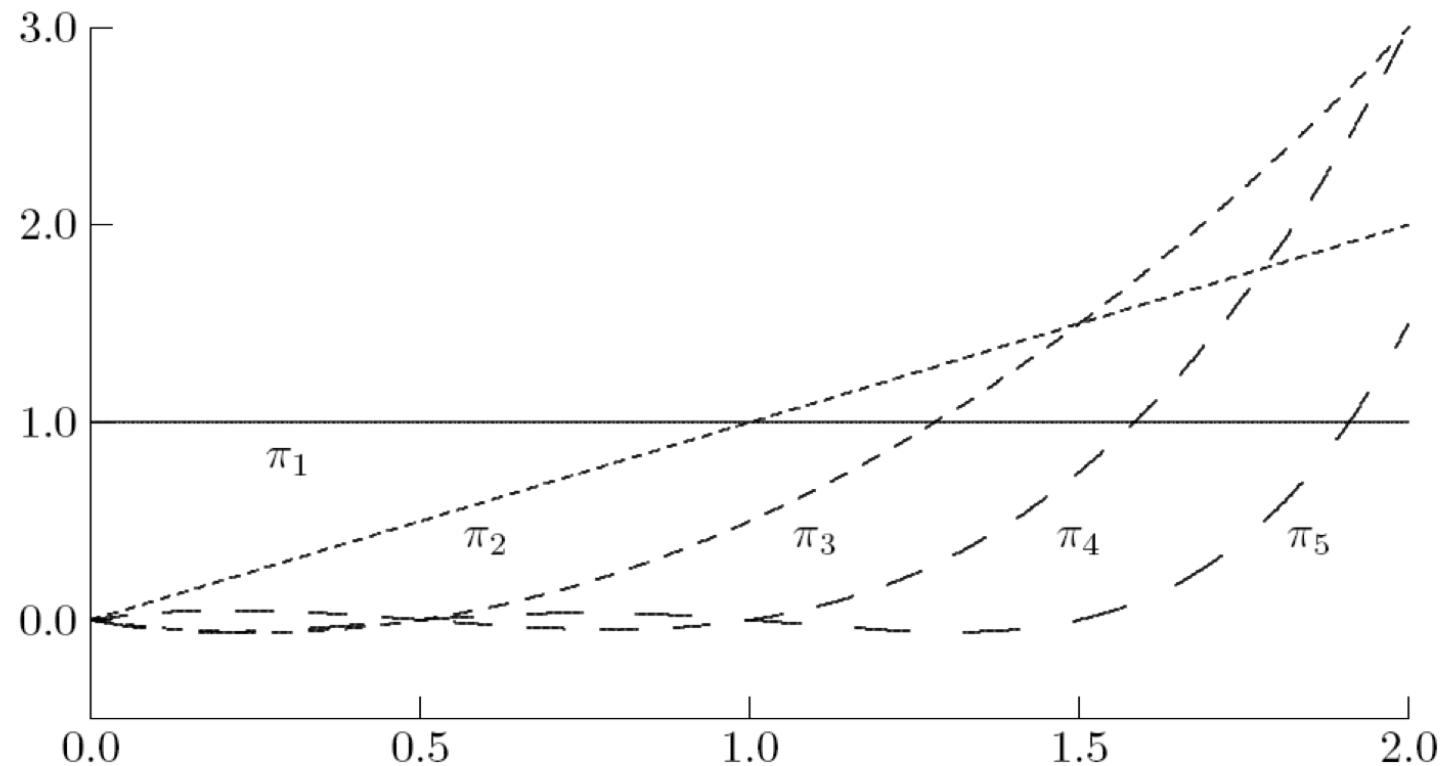
$$\pi_j(t) = \prod_{k=1}^{j-1} (t - t_k), \quad j = 1, \dots, n$$

- Newton interpolating polynomial has form

$$\begin{aligned} p_{n-1}(t) &= x_1 + x_2(t - t_1) + x_3(t - t_1)(t - t_2) + \\ &\quad \cdots + x_n(t - t_1)(t - t_2) \cdots (t - t_{n-1}) \end{aligned}$$

- For  $i < j$ ,  $\pi_j(t_i) = 0$ , so basis matrix  $A$  is lower triangular, where  $a_{ij} = \pi_j(t_i)$ .

# Newton basis functions

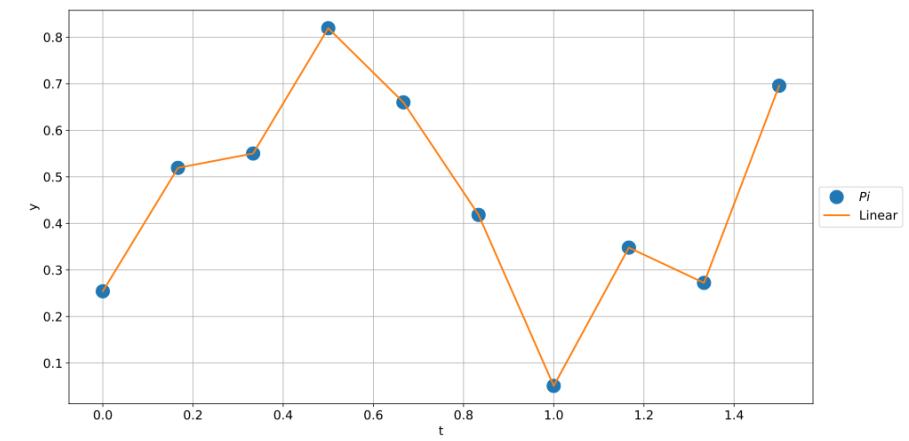
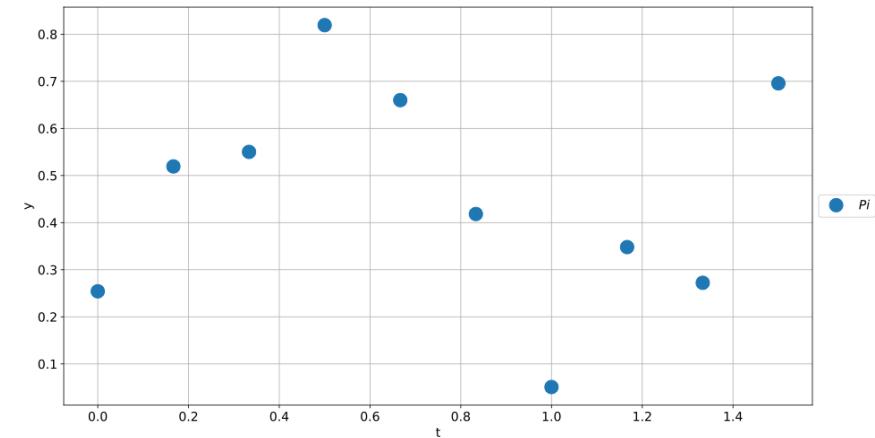


# Piecewise polynomial interpolation

- Motivation:
  - Fitting single polynomial to large number of data points is likely to yield unsatisfactory behavior in interpolant
- Main advantage:
  - Large number of data points can be fit with low-degree polynomials
- How:
  - Given data points  $(t_i, y_i)$ , different function is used in each subinterval  $[t_i, t_{i+1}]$ 
    - $t_i$  is called knot or breakpoint, at which interpolant changes from one function to another

# Piecewise polynomial interpolation

- Discussion: Could you provide an example of a piecewise polynomial interpolation?
- Simplest example is piecewise linear interpolation, in which successive pairs of data points are connected by straight lines
  - Discussion: what are the drawbacks of linear interpolation?

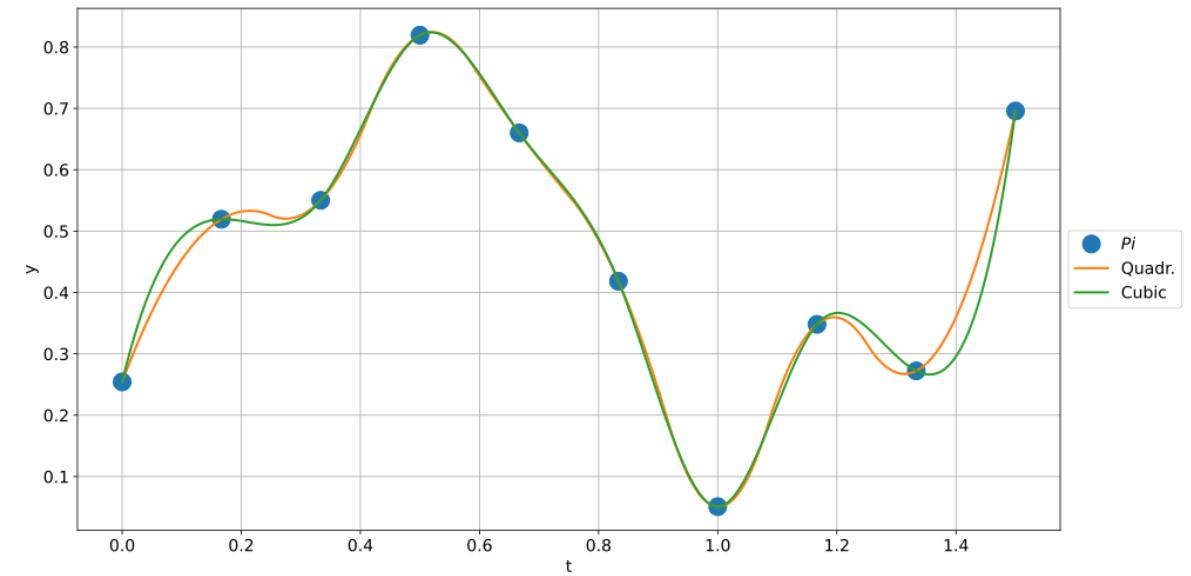
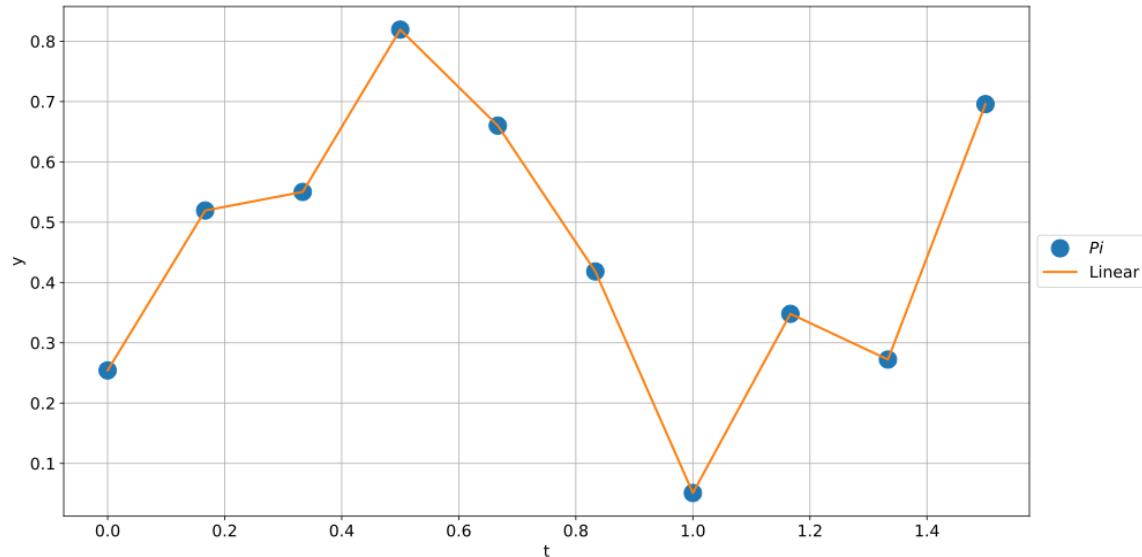


# Spline interpolation

- A spline is a smooth piecewise polynomial function.
  - Two popular model:
    - Quadratic spline, Cubic spline
- **Quadratic** spline interpolation
  - each segment is a **second-degree polynomial** function.
  - Formally, we have data points  $(t_i, y_i), i = 1, \dots, n$
  - For each interval  $[t_i, t_{i+1}]$ , we define a quadratic polynomial
    - $f_i(t) = a_i + b_i(t - t_i) + c_i(t - t_i)^2$ .
    - There are  $n - 1$  such polynomials (one per interval).
  - Discussion: how many coefficients need to be determined? How many equations do we need?
    - $3(n - 1)$

# Illustration of piecewise polynomial interpolation (scipy.interpolate)

- Piecewise linear
- Spline – quadratic
- Spline - cubit



# Summary

- Interpolating function fits given data points **exactly**, which is not appropriate if data are noisy
- Interpolating function given by **linear combination of basis functions**, whose coefficients are to be determined
- Existence and uniqueness of interpolant depend on whether **number of parameters** to be determined matches **number of data points** to be fit
- Piecewise polynomial (e.g., spline) interpolation can fit **large number of data points** with low-degree polynomials
- Cubic spline interpolation is excellent choice when **smoothness** is important

# Finally, ...

- HW4 due tonight
- Final practice exam and solution reviews next week
- Looking forward to your final presentations on Mon Dec 8!
- I'll teach CSI 436 Machine Learning next Spring.
  - If you enjoy my teaching or want to learn more about AI/ML, feel free to register
  - Let's see how numerical methods are applied in real-world exciting algorithms and applications!